

UMD DATA605 - Big Data Systems

Course Intro

Big Data

Data Models

Dr. GP Saggese

gsaggese@umd.edu

with thanks to Prof.

Alan Sussman

Amol Deshpande

UMD DATA605 - Big Data Systems

Course Intro

Big Data

Data Models

Goals

- Learn to model and reason about data
- Learn how to process and manipulate data in different ways
- Introduce a variety of data models
 - Decide which data model is appropriate for different applications
- Learn to use a variety of data management systems
 - Decide which system is appropriate for a data management scenario
- Learn how to build data processing pipelines
- Learn how to build a big-data system end-to-end
 - Class project!
 - Team work

Todos

- [DATA605 - ELMS / Canvas site](#)
 - Make sure to enable notifications
 - How to get in touch with me / TA
- [DATA605 - Schedule](#)
- Clone [DATA605 - GitHub repo](#)
- Begin setup of your computing environment
 - Install Docker on your laptop
 - Instructions are in the class repo

Tools We Will Learn To Use

- **Languages:** mainly Python
- **Dev tools**
 - Bash / Linux OS
 - Git (understand data model, branching)
 - GitHub (PRs, work with issues)
 - Code editors (PyCharm, Visual Studio, vim / emacs)
 - Data science libs (pydata stack: numpy, Pandas, sklearn, matplotlib)
 - Jupyter notebooks
 - Docker
 - Unit testing framework
- **Big data tools**
 - ETL pipelines
 - Relational DBs (PostgreSQL)
 - NoSQL DBs (HBase, MongoDB, Couchbase, Redis)
 - Graph DBs (Neo4j, OrientDB, AllegroGraph, GraphX, Giraph)
 - Computing framework (Hadoop, Spark, Dask, Storm, Spark Streaming)
 - Workflow manager (Airflow)
 - Cloud Services (AWS)
- **Tutorials** for all tools we use for the class project

Soft Skills to Succeed in the Workplace

- **Goal:** model the class project to prepare you for the workplace
- Skills that people out from college often don't have
 - Work in a team
 - Design software architecture (e.g., OOP, Agile, Design Patterns)
 - Comment your own code
 - Write external documentation (e.g., tutorials, manuals, how-tos)
 - Write code so that other people can understand (including future-you)
 - Read other people's code
 - Follow code conventions (e.g., PEP8, Google Code)
 - Be clear in communications (e.g., in emails, Slack)
 - File a bug report
 - Reproduce (aka repro) a bug
 - Have a sense of CS constants
 - Have a sense of how an OS works (e.g., virtual memory, processes)

Class Project

- [DATA605 - Class project](#)
 - Build an end-to-end big data system importing data, computing, orchestrating the tasks
 - Teams of 4-5 students
 - GP and the TA will create well-balanced teams
 - Review 2 examples of complete projects
 - 5 deliverables with deadlines graded individually (20% of final grade)
 - Every class ~1 hour of group work with GP + TA
 - Peer evaluation of the project
 - Grad and undergrad students will be assigned final grades separately
 - Will re-evaluate based on how projects go
- [Student survey](#)
 - Collect some information about you to organize groups for the class project
- Check out an example of Big Data System
 - [Sorrentum project](#) and [Git repo](#)

Why I Am Interested in Big Data

- <https://www.linkedin.com/in/gpsaggese/>
 - Feel free to connect on LinkedIn so we can stay in touch
- Open-source project I've started
 - [Sorrentum](#)
 - [GitHub](#)

Your Turn

- Introduce yourself
- Tell us something about yourself
 - E.g., hobby, peculiar habits, ...

UMD DATA605 - Big Data Systems

Course Intro

Big Data

Data Models

Data Science

- **Promises of data science** (DS)
 - Give a competitive advantages
 - Make better strategic and tactical business decisions
 - Optimize business processes
 - Detect strategic blind spots
- **Data science is not new**, it was called:
 - Decision support
 - Business intelligence
 - Predictive analytics
 - Operation research
 - ...
- What has changed is that **learning and applying DS is easy**
 - No need for consulting groups
 - Tools are open-source
 - E.g., Python + pydata stack (numpy, scipy, Pandas, sklearn)
 - Large data sets
 - Cheap computing (e.g., AWS, Google Cloud)

Motivation: Data Overload

- McKinsey, 2013: “data science is the **number one catalyst for economic growth**”
 - Growing amount of data
 - Cloud computing
 - Build machine models to make better business decisions
- **Explosion of data in every domain**
 - Sensing devices / networks monitor processes 24/7
 - E.g., temperature of your room, your vital signs, pollution in the air
 - Increasingly sophisticated smart-phones (6.5b smart phones, 80% of the world population)
 - Internet and social networks make it easy to publish data
 - Internet of Things (IoT), everything is connected to the internet
 - E.g., power supply, toasters
 - Scientific experiments and simulations produce enormous volumes of data
 - Datafication: turn all aspects of life into data
 - E.g., what you like/enjoy turned into a stream of your "likes"
- **Challenges**
 - How to handle increasing amount data?
 - How to extract interesting actionable insights and scientific knowledge?

Scale of Data Size

- Megabyte = 2^{10} \approx 10^6 bytes
 - Typical English book
- Gigabyte = 10^9 bytes
 - 1/2 hour of video
 - Wikipedia (compressed, without media): 22GB
- Terabyte = 1m MB
 - Human genome: ~1TB
 - 200,000 photos
 - LHC generates 100TB of data per day
 - \$50 to buy 1TB HDD, \$23 / mo on AWS S3
- Petabyte = 1000 TB
 - 13 years of HD video
 - 65 copies of Library of Congress
 - \$250k / year on AWS S3
- Exabyte = 1m TB
 - Global yearly Internet traffic in 2004
- Zetabyte = 1b TB = 10^{21} bytes
 - Global yearly Internet traffic in 2016
- Yottabytes = 10^{24} bytes
 - 1TB today costs ~\$100, Yottabyte costs \$100t
- Brontobytes = 10^{27} bytes

How big is a Yottabyte?

TERABYTE

Will fit 200,000 photos or mp3 songs on a single 1 terabyte hard drive.



PETABYTE

Will fit on 16 Backblaze storage pods racked in two datacenter cabinets.



EXABYTE

Will fit in 2,000 cabinets and fill a 4 story datacenter that takes up a city block.



ZETTABYTE

Will fill 1,000 datacenters or about 20% of Manhattan, New York.



YOTTABYTE

Will fill the states of Delaware and Rhode Island with a million datacenters.



The Cost

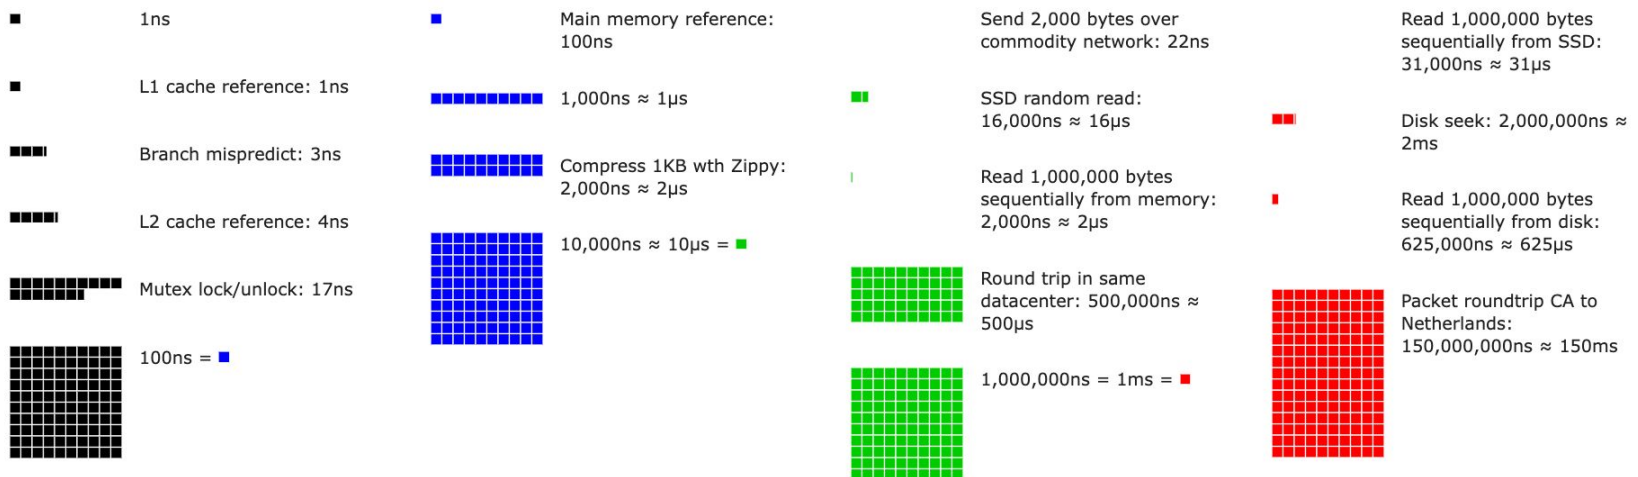
The cost of buying a 1 terabyte hard drive today is \$100. It would cost \$100 Trillion dollars to buy a yottabyte of storage for just the hard drives.



Constants that everybody should know

From [Latency Numbers Every Programmer Should Know \(by year\)](#)

- A CPU running at 3GHz executes an instruction every 0.3ns
- L1 cache reference / register: 1ns
- L2 cache reference: 4ns
- Main memory reference: 100ns
- Send 1KB over network: 10ns
- Read 1MB from memory: 2us
- SSD random read: 16us
- Disk seek: 2ms
- Packet round-trip from CA to Netherland: 150ms



Big Data Applications

- **Personalized marketing**

- Target each consumer instead of the consumers at large
- E.g., Amazon personalizes suggestions using signals from:
 - Your shopping history
 - What you have searched for (or clicked, watched)
 - Other consumers and trends
 - Reviews (through NLP and sentiment analysis)
- Brands want to understand how customers relate to products
 - Use sentiment analysis from
 - social media
 - on-line reviews
 - blogs
 - surveys
 - Positive, negative, neutral feeling
- E.g.,
 - In 2022, \$600b spent on digital marketing
 - [50 Stats Showing The Power Of Personalization](#)

Big Data Applications

- **Mobile advertisement**

- Mobile phones are ubiquitous
 - 80% of world population has one
- Integration of on-line and off-line databases, e.g.,
 - GPS location
 - Search history, credit history, credit card transactions
- E.g.,
 - You've bought a new house
 - You google questions about house renovations
 - You watch shows about renovations
 - Your phone tracks where you are
 - Google sends you coupons for the closest Home Depot

Big Data Applications

- **Biomedical data**

- Personalized medicine
 - Patients can receive treatment tailored to them to maximize efficacy
 - Genetics, daily activities, environment, habits
- Genome sequencing

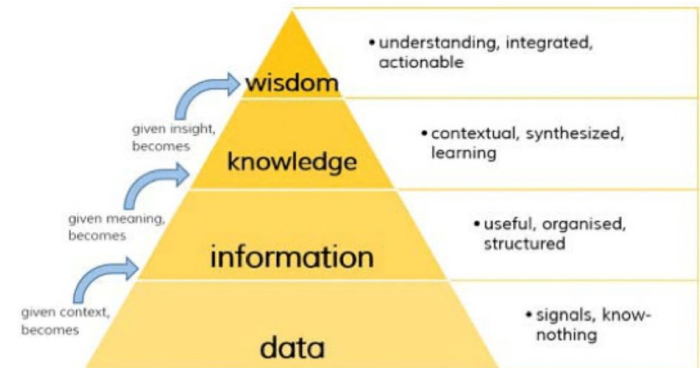
Big Data Applications

- **Smart cities**

- Interconnected mesh of sensors
 - E.g., traffic sensors, camera networks, satellites
- Goals:
 - Monitor air pollution
 - Minimize traffic congestion
 - Optimal urban services
 - Maximize energy savings


Goal of Data Science

- Goal of data science: from data to wisdom
 - Data (raw bytes)
 - Information (organized, structured)
 - Knowledge (learning)
 - Wisdom (actionable, understanding)
- Combine streams of big data to generate new data
 - New data can be “big data” itself
- Insights enable decisions and actions



Four V's of Big Data

- Big data often exhibits certain characteristics and qualities
- **Volume**
 - Vast amount of data is generated
- **Variety**
 - Different forms
 - E.g., text, images, voice, geospatial data
- **Velocity**
 - Speed at which data is generated
- **Veracity**
 - Biases, noise, abnormality in data
 - Uncertainty, trustworthiness
- (**Valence**)
 - Connectedness of big data in the form of graphs
- (**Value**)
 - Challenges of big data can distract from focusing on how big data benefits an organization



Four V's of Big Data

- **Volume**

- Exponentially increasing amount of data
- Every day 2.5 exabytes (1m of TB) of data is generated
 - 90% of all the data in the world was generated in the last 2 years
 - Total amount of stored data doubles every 1.2 years
- Twitter: 500M tweets / day (2022)
- Google processes 8.5b queries / day (2022)
- Facebook generates 4PB of data / day (2022)
- Walmart: 2.5PB of unstructured data / hour (2022)

- **Variety**

- Different form
 - Structured data (e.g., spreadsheets, relational data)
 - Semi-structured data (e.g., natural language text, sales receipts, your class notes)
 - Unstructured data (e.g., photos, videos)
- Different formats (e.g., flat files, CSV, XML, JSON)

Four V's of Big Data

- **Velocity**

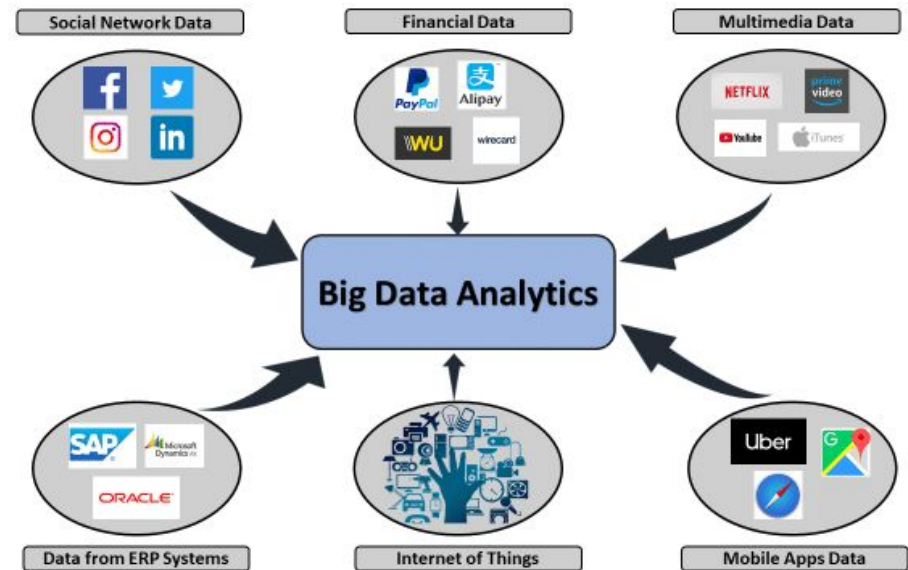
- Relates to the speed at which data is generated
 - E.g., sensors generating data streams
- Sometimes data can be processed off-line
- Real-time analytics requires data to be consumed as fast as it is generated

- **Veracity**

- Relates to data quality
- How to remove noise and bad data?
- How to fill in missing values?
- What is an outlier?
- How do you decide what data to trust?

Sources of Big Data

- We can distinguish Big Data in terms of its source
 - **Machines**
 - **People**
 - **Organizations**



Sources of Big Data: Machines

- **Machines** generate data
 - Real-time sensors (e.g., sensors on Boeing 787)
 - Cars
 - Website tracking
 - Personal health trackers
 - Scientific experiments
- Pros
 - Highly structured
- Cons
 - Can't be easily moved, but need to be computed in-situ or in centralized fashion
 - Streaming, not batch

Sources of Big Data: People

- **People** and their activities generate data
 - Social media (e.g., Facebook, Instagram, Twitter, LinkedIn)
 - Video sharing (e.g., YouTube, TikTok)
 - Blogging and commenting on a website
 - Internet searches
 - Text messages (e.g., SMS, Whatsapp, Signal, Telegram)
 - Personal documents (e.g., Google Docs, emails)
- Pros
 - Highly valuable for business intelligence
 - Allow personalization
- Cons
 - Typically unstructured data
 - Text, images
 - It takes investment before you can reap the value
 - Acquire → Store → Clean → Retrieve → Process → Insights
 - Surveillance capitalism

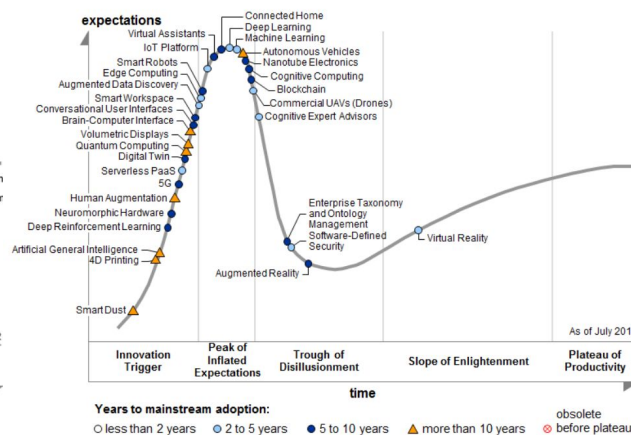
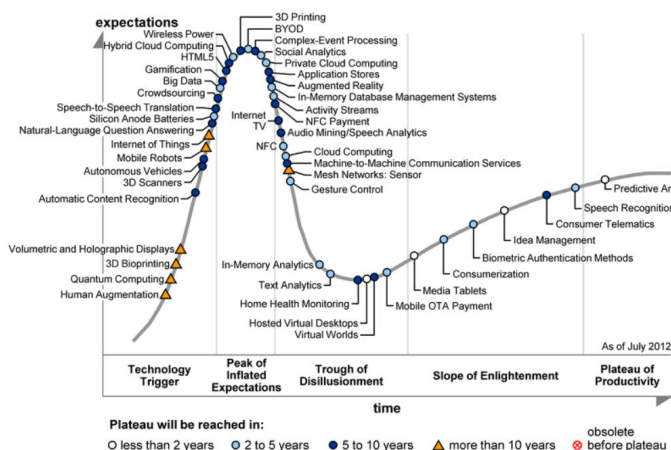



Sources of Big Data: Orgs


- **Organizations** generate data
 - Commercial transactions
 - Credit cards
 - E-commerce
 - Banking
 - Financial information
 - Medical records
 - Transactions
 - Clicks on a website
- Pros
 - Highly structured
- Cons
 - Store every event in the past to predict the future
 - Stored in “data silos” with different data models
 - Each department has its own data system
 - Additional complexity, missing opportunities, data is outdated / not visible
 - Cloud computing helps (e.g., data lakes, data warehouses)

Big Data + Data Science to the Rescue

- *Big data and data science* often used synonymously
 - Aka data analytics, data mining, business intelligence, predictive analytics
 - Loosely used for any process where interesting information are inferred from data
- Data scientist called the “sexiest job” of the 21st century
 - The term has becoming very muddled at this point



Hype Cycle for Emerging Tech, 2022



Is Data Science Just Hype?

- **No**

- Extracting insights and knowledge from data:
 - Is very important
 - Will continue to increase in importance
- Big data techniques are revolutionizing the world in many domains
 - E.g., education, food supply, disease epidemics, ...

- **But**

- Not much different from what statisticians have been doing for many years

- **What is different? What has changed?**

- Much more data is digitally available than ever before
- Easy-to-use programming frameworks (e.g., Hadoop, Spark) = much easier to analyze it
- Cloud computing
- Often large-scale data + simple algorithms > small data + complex algorithms

Key Shifts Before / After Big-Data

- 1) Datasets: small, curated, clean → large, uncurated, messy
 - Before:
 - Statistics based on small, carefully collected random samples
 - Costly and careful planning for experiments
 - Hard to do fine-grained analysis
 - Today:
 - Easily collect huge volume of data
 - Feed it into algorithms
 - Usually the signal is strong enough to overcome the noise
 - 2) Causation → Correlation
 - Goal: figure out what caused what
 - Causation very hard to figure out
 - Today: give up causation for correlation
 - Find out that two things are correlated and often that's enough
 - 3) "Datafication" = process of converting abstract things into concrete data
 - E.g., "sitting posture" is datafied by capturing information with 100's of sensors placed in your seat
 - Your preference is datafied into a stream of your likes
- From: [Rise of Big Data](#), 2013

Examples: Election Prediction

- [Nate Silver and the 2012 Elections](#)
 - 49 out of 50 in calling each state in 2008 US elections
 - 50 out of 50 in 2012 US elections
 - Didn't work that well in 2016
- [Some reasons why he got things right](#)
 - Many sources of data, irrespective of quality
 - Incorporation of historical accuracy
 - Use of statistical models
 - Understanding correlations
 - Monte-Carlo simulations to compute the probabilities of electoral college
 - Focus on probabilities instead of predictions
 - Great communication and presentation skills




Examples: Google Flu Trends

- Google Flu Trends
 - Early warning of flu outbreaks by analyzing search queries
 - What terms people searched for (45 search terms used to create a single model)
 - IP to determine location
 - Application of “collective intelligence”
 - Predict regional outbreaks of flu up to 1 or 2 weeks ahead of CDC
 - 5% to 20% of US population contract flu every year and 40k deaths
 - Earlier warnings allow prevention and control
 - Service in activity from 2008 to 2015
- Caveat: accuracy not as good any more
 - Google claimed 97% accuracy
 - Out of sample accuracy lower (overshot CDC data by 30%)
 - People search about flu without knowing how to diagnose flu (e.g., people searching for “fever” and “cough”)
 - [Google Flu Trends: The Limits of Big Data](#)

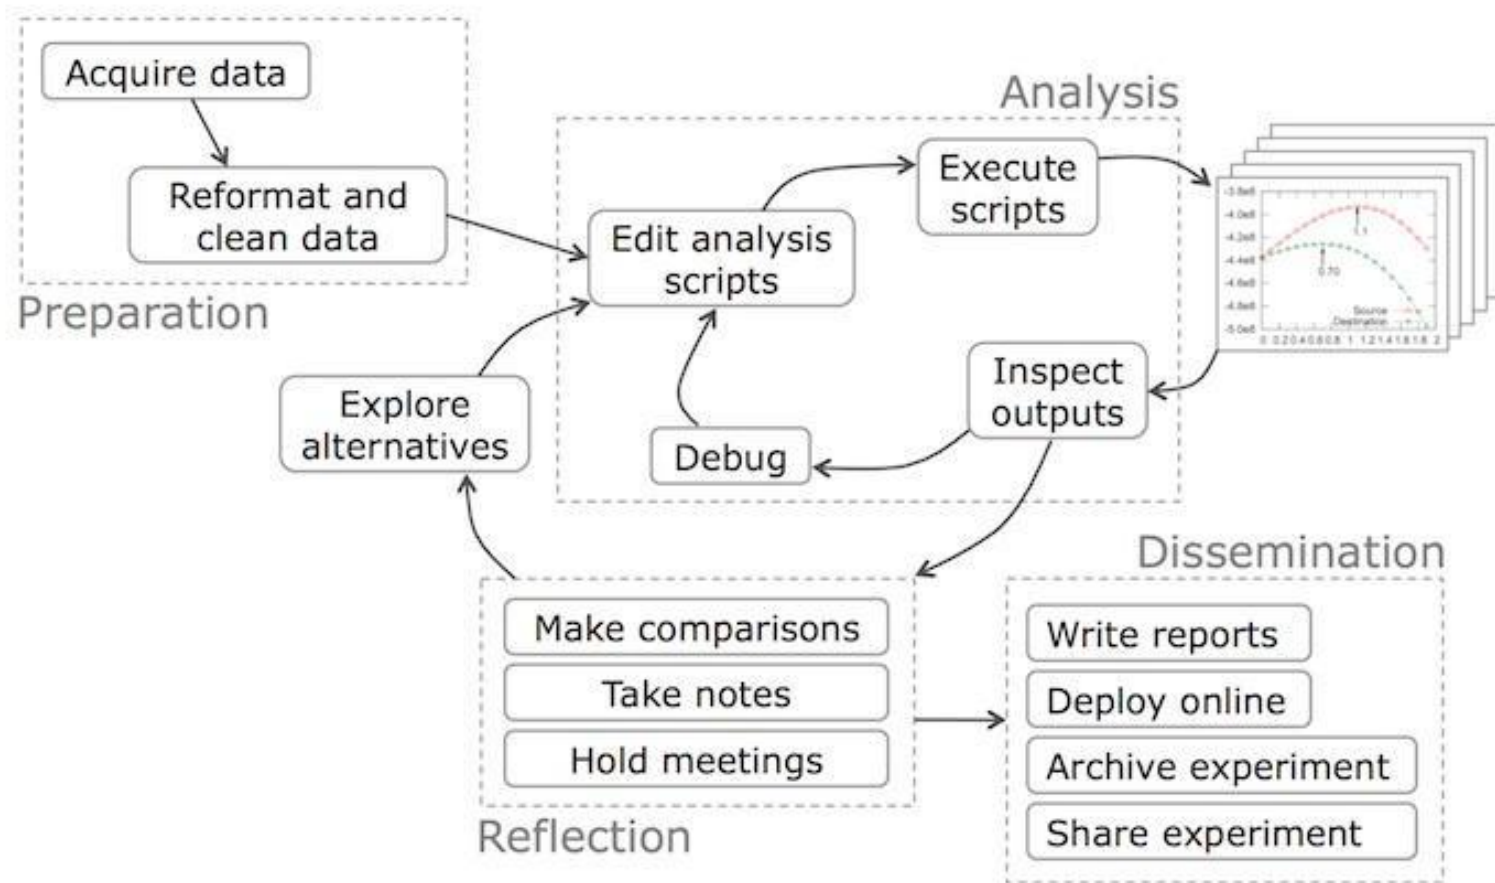
Data Scientist

- Very ambiguous and ill-defined term



- From [Drew Conway's Venn Diagram](#)

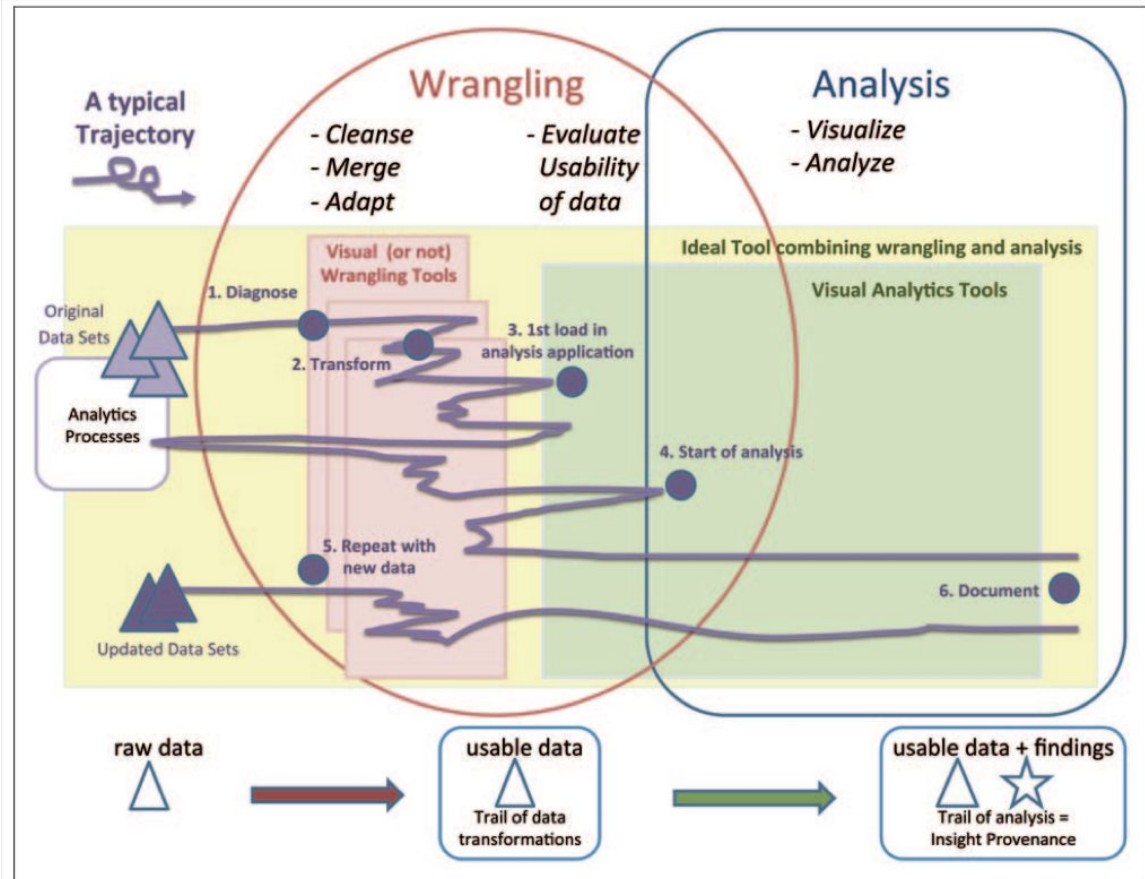
Typical Data Science Workflow



- From [Data Science Workflow](#)

Where Data Scientist Spends Most Time

- 'Janitor Work' in Data Science
- Research Directions in Data Wrangling
 - Estimates that 80-90% of the work is in data cleaning and wrangling



What a Data Scientist Should Know

- From: [How to hire a data scientist](#)
- Data grappling skills
 - How to move data around and manipulate it with some programming language
 - Scripting languages (e.g., Python)
 - Data storage tools like relational databases, key-value stores
 - Programming frameworks like SQL, Hadoop, Spark, etc.
- Data visualization experience
 - How to draw informative pictures of data
 - Many tools (e.g., D3.js, plotting libraries)
 - Harder question is knowing what to draw
- Knowledge of statistics
 - E.g., error-bars, confidence intervals
 - Python libraries; Matlab; R
- Experience with forecasting and prediction
 - Basic machine learning techniques
- Communication skills
 - Communicate the findings

We focus here!!!

UMD DATA605 - Big Data Systems

Course Intro

Big data

Data Models

Data Models

- **Data modeling**

- = process of representing and capturing the structure and properties of real-world entities
- Process of abstraction: real-world → representation

- **Data model**

- = description of how data is represented (e.g., relational, key-value) and accessed (e.g., allowed operations, how to query)
- E.g., schema in a DB describe a specific collection of data, using a given data model

- **Why do we need data model?**

- Need to know the structure of the data (to some extent) to be able to write general purpose code
- Allow to share data across programs, organizations, systems
- Need to integrate information from multiple sources
- Preprocess data to make access efficient (e.g., building an index on a data field)

Multiple Layers of Data Modeling

- **Physical layer**


- How is the data physically stored
- How to represent complex data structures (e.g., B-trees for indexing)

- **Logical layer**

- Type of information stored
- Entities
- Entity attributes
- Relationships among the above

- **Views**

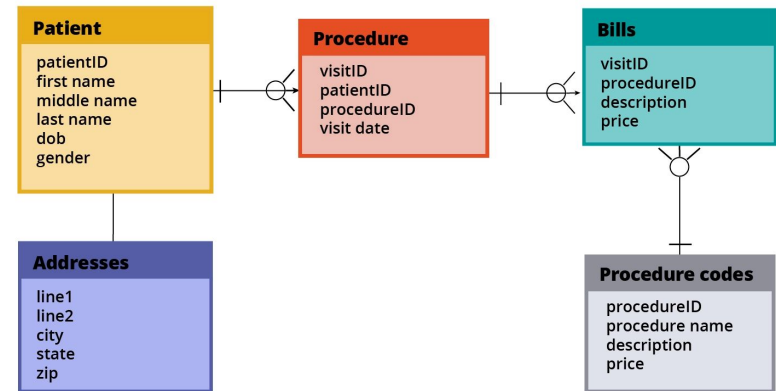
- Restrict information flow
- Security and/or ease-of-use



Data Models: Logical Layer

- **Modeling constructs**

- A collection of concepts used to represent the structure in the data, e.g.,
 - Types of entities
 - Entity attributes
 - Types of relationships between entities
 - Types of relationships between attributes



- **Integrity constraints**

- Constraints to ensure data integrity
 - Goal: avoid errors and data inconsistencies
 - E.g., a field can't be empty, is an integer

- **Manipulation constructs**

- E.g., insert, query, update, delete data

Examples of Data Models

- We will cover:
 - Relational model (SQL)
 - Entity-relationship (ER) model
 - XML
 - Object-oriented (OO)
 - Object-relational
 - RDF
 - Property graph
- Serialization formats are also data models
 - CSV
 - Parquet
 - JSON
 - Protocol Buffer
 - Avro / Thrift
 - Python Pickle

Good Data Models

- We would like a data model to be:
 - Expressive
 - Capture real-world data well
 - Easy to use
 - Good performance
- Tension between the above characteristics
 - More powerful models
 - Can represent more datasets
 - Harder to use, to query
 - Less efficient
- The evolution of data modeling tools is an attempt to capture the structure in the data
 - Structured data → Relational DBs
 - Semi-structured web data → XML
 - Unstructured data → NoSQL DBs

Data Independence

- **Logical data independence**
 - Can change the representation of data without changing programs that operate on it
- **Physical data independence**
 - Can change the layout of data on disk and programs won't change
 - Index the data
 - Partition / distribute / replicate the data
 - Compress the data
 - Sort the data

Databases: A Brief History

- **1960s: Early beginning**

- Computers finally become attractive technology
- Enterprises start adopting computers
- Most applications initially used their own data stores
 - Each application had its own format
 - Data was basically unavailable to other programs
- **Database:** term coined in military information systems to denote "shared data banks" by multiple applications
 - Define a data format
 - Store it as a "data dictionary" (schema)
 - Implement general-purpose "database management" software to access data
- Issues:
 - How to write data dictionaries?
 - How to access data?
 - Disadvantages of integration: integrity, security, privacy concerns
 - Who controls the data?

Databases: A Brief History

- Birth of "hierarchical model" and "network model"
 - Both allowed connecting records of different types
 - E.g., connect accounts with customers
 - Network model attempted to be very general and flexible
- IBM designed [IMS hierarchical database](#) in 1966 for the Apollo space program
 - Predates *hard disks*
 - Still around today!
 - .. *more than 95 percent of the top Fortune 1000 companies use IMS to process 50 billion transactions a day and manage 15 million gigabytes of critical business data* (from [IBM Website on IMS](#))
- Cons:
 - Hierarchical / network models exposed too much of the internal data (e.g., structures / pointers to the users)
 - Leaky abstraction

Relational, Hierarchical, Network model

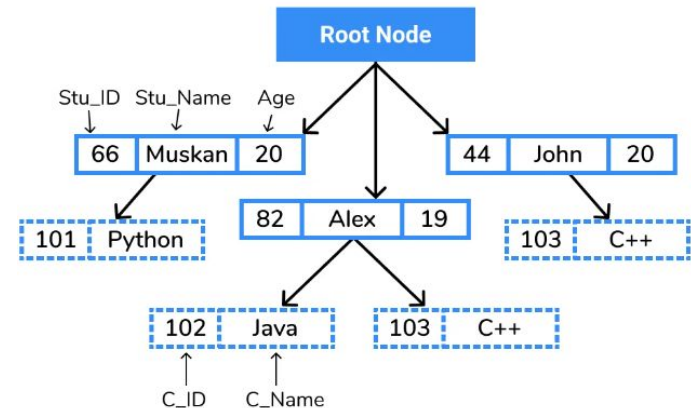
- Relational model

- Data is represented as tuples grouped in relation
- Omnipresent SQL

Customer ID	Tax ID	Name	Address	[More fields...]
1234567890	555-5512222	Ramesh	323 Southern Avenue	...
2223344556	555-5523232	Adam	1200 Main Street	...
3334445563	555-5533323	Shweta	871 Rani Jhansi Road	...
4232342432	555-5325523	Sarfaraz	123 Maulana Azad Sarani	...


- Hierarchical model

- Data is organized into a tree-like structure
 - Each record has one parent record and many children
 - Records connected through links
- Resurgence in 1990s with XML DBs



- Network model

- Data is organized in a graph
 - Each record can have multiple parent and child records
- Resurgence in 2010s with graph DBs



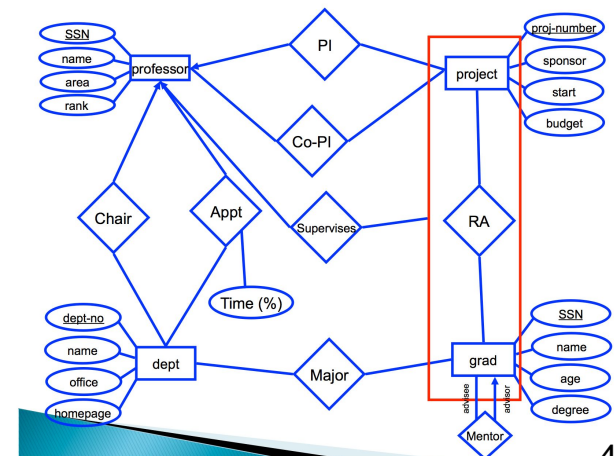

Databases: A Brief History

- **1970s: Relational model**

- Set Theory, First-order predicate logic
 - Ted Codd developed the Relational Model
- Elegant, formal model
 - Provided almost complete *data independence*
 - Users didn't need to worry about how the data was stored, processed
- High level query language
 - SQL based on relational algebra
- Notion of *normal forms*
 - Allowed one to reason about
 - Remove redundancies
- Influential projects: INGRES (UC Berkeley), System R (IBM)
 - Didn't care about IMS/IDMS compatibility (as IBM had to)
- Many debates between Relational Model vs Network Model proponents

Entity-Relationship Model

- 1976: Peter Chen proposed "Entity-Relationship Model"
- Data model describing knowledge in terms of entities and relationships
- **Entities** are physical or logical objects that can be uniquely identified
 - “Nouns”
- **Relationships** between entities
 - “Verbs”
- An ER model can be mapped onto a relational DB
 - Entities, relationships -> tables



Databases: A Brief History

- **1980s: Widespread acceptance of relational model**
 - SQL emerged as a standard, in large part because of IBM's backing
 - Enriching the expressive power of relational model
 - Set-valued attributes, aggregation, etc.
- **Late 80's**
 - Object-oriented DB
 - Store objects instead of tables
 - Get around *impedance mismatch* between programming languages and databases
 - Object-relational DB
 - Allow user-defined types
 - Get many benefits of object-oriented while keeping the essence of relational model
 - No expressive difference from pure relational model

Object-Oriented

- OOP is a data model
 - Object behavior is described through data (stored as fields) and code (in the form of methods)
- Composition
 - Aka `has-a` relationships
 - E.g., an Employee class has an Address class
- Inheritance
 - Aka `is-a` relationships
 - E.g., an Employee class derives from a Person class
- Polymorphism
 - Code executed depends on the class of the object
 - One interface, many implementations
 - E.g., `draw()` method on a Circle vs Square object, both descending from Shape class
- Encapsulation
 - E.g., private vs public fields / members
 - Prevents external code from being concerned with inner workings of an object

Databases: A Brief History

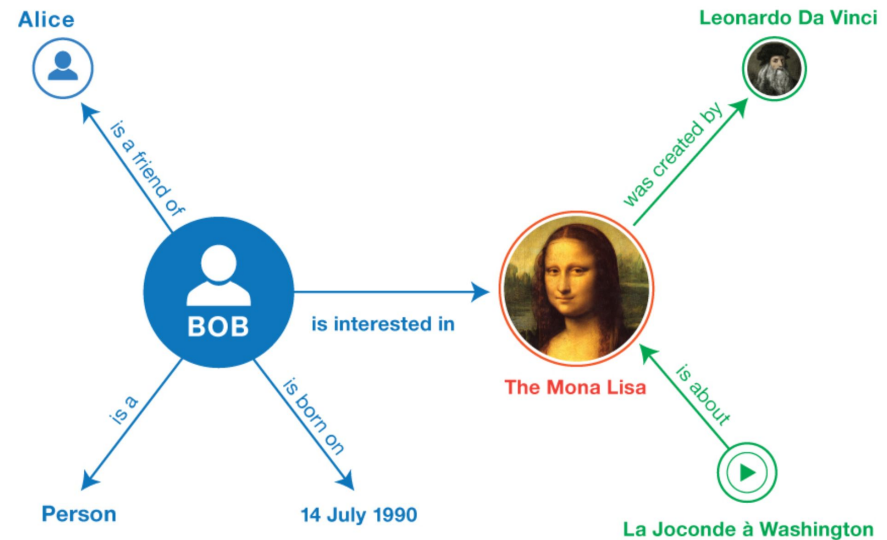
- **Late 90's-today**
- Web emerges
- XML: eXtensible Markup Language
 - Intended for *semi-structured* data
 - Tree-like structure
 - Flexible schema

```
<?xml version="1.0" encoding="UTF-8"?>
<!-- Edited by XMLSpy -->
<CATALOG>
  <CD>
    <TITLE>Empire Burlesque</TITLE>
    <ARTIST>Bob Dylan</ARTIST>
    <COUNTRY>USA</COUNTRY>
    <COMPANY>Columbia</COMPANY>
    <PRICE>10.90</PRICE>
    <YEAR>1985</YEAR>
  </CD>
  <CD>
    <TITLE>Hide your heart</TITLE>
    <ARTIST>Bonnie Tyler</ARTIST>
    <COUNTRY>UK</COUNTRY>
    <COMPANY>CBS Records</COMPANY>
    <PRICE>9.90</PRICE>
    <YEAR>1988</YEAR>
  </CD>
  ...

```

Resource Description Framework

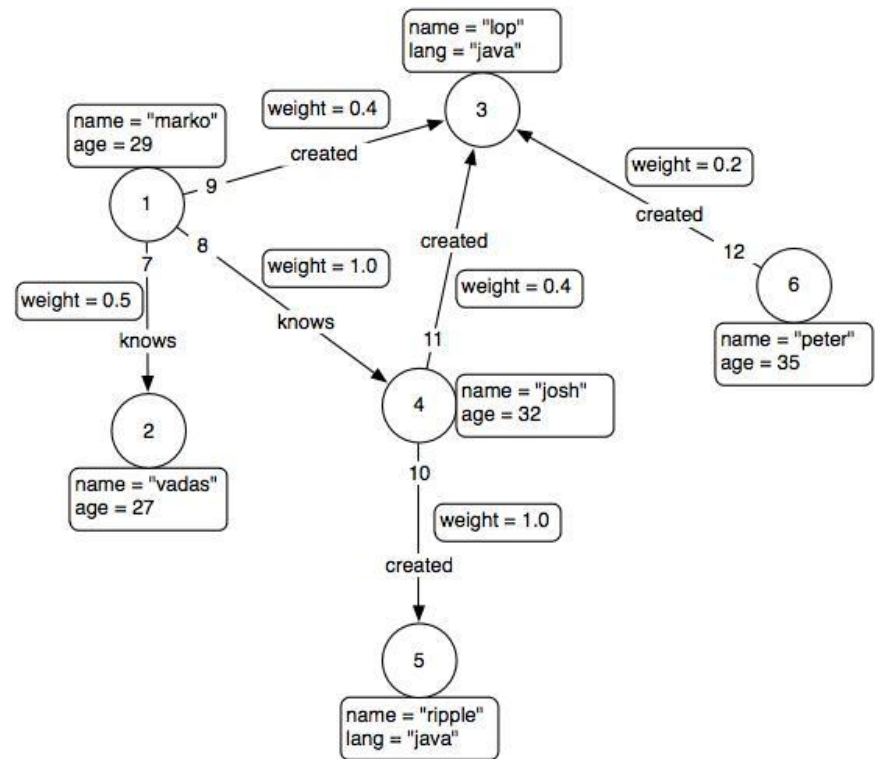
- Aka RDF
- Key construct: a "subject-predicate-object" triple
 - E.g.,
 - subject=sky
 - predicate=has-the-color
 - object=blue
- Can be mapped to a labeled, directed multi-graph
 - More general than a tree
- Typically stored in:
 - Relational DBs
 - Dedicated "triple-stores" DBs



```
01 <http://example.org/bob#me>
   <http://www.w3.org/1999/02/22-rdf-syntax-ns#type>
   <http://xmlns.com/foaf/0.1/Person> .
02 <http://example.org/bob#me> <http://xmlns.com/foaf/0.1/knows>
   <http://example.org/alice#me> .
03 <http://example.org/bob#me> <http://schema.org/birthDate>
   "1990-07-04"^^<http://www.w3.org/2001/XMLSchema#date> .
04 <http://example.org/bob#me>
   <http://xmlns.com/foaf/0.1/topic_interest>
   <http://www.wikidata.org/entity/Q12418> .
05 <http://www.wikidata.org/entity/Q12418>
   <http://purl.org/dc/terms/title> "Mona Lisa" .
06 <http://www.wikidata.org/entity/Q12418>
   <http://purl.org/dc/terms/creator>
   <http://dbpedia.org/resource/Leonardo_da_Vinci> .
07
   <http://data.europeana.eu/item/04802/243FA8618938F4117025F17A8B81
   3C5F9AA4D619> <http://purl.org/dc/terms/subject>
   <http://www.wikidata.org/entity/Q12418> .
```

Property Graph Model

- Graph:
 - with vertices and edges
 - with properties associated with each edge and vertex
- Typically stored in:
 - Relational DBs
 - Graph DBs



Serialization Formats

- Programs need to send data to each other (on the network, on disk)
 - E.g., Remote Procedure Calls
 - Several recent technologies based around schemas
 - JSON, YAML, Protocol Buffer
- Serialization formats are data models

JSON

- JSON = JavaScript Object Notation
- Data is nested dictionaries and arrays
- Very similar to XML
 - More human-readable
 - Less boilerplate
 - Executable in JavaScript (and Python)

```
{
  "firstName": "John",
  "lastName": "Smith",
  "isAlive": true,
  "age": 25,
  "height_cm": 167.6,
  "address": {
    "streetAddress": "21 2nd Street",
    "city": "New York",
    "state": "NY",
    "postalCode": "10021-3100"
  },
  "phoneNumbers": [
    {
      "type": "home",
      "number": "212 555-1234"
    },
    {
      "type": "office",
      "number": "646 555-4567"
    }
  ],
  "children": [],
  "spouse": null
}
```

Protocol Buffers

- Developed by Google
- [Open-source](#)
- Represent data structures in a
 - Language agnostic
 - Platform agnostic
 - Versioning
- Schema is mostly relational
 - Optional fields
 - Types
 - Default values
 - Structures
 - Arrays
- Schema specified using a .proto file
- Compiled by protoc to produce C++, Java, or Python code to initialize, read, serialize objects

```
message Person {
  optional string name = 1;
  optional int32 id = 2;
  optional string email = 3;

  enum PhoneType {
    MOBILE = 0;
    HOME = 1;
    WORK = 2;
  }

  message PhoneNumber {
    optional string number = 1;
    optional PhoneType type = 2;
  }

  repeated PhoneNumber phones = 4;
}
```

```
import addressbook_pb2
person = addressbook_pb2.Person()
person.id = 1234
person.name = "John Doe"
person.email = "jdoe@example.com"
phone = person.phones.add()
phone.number = "555-4321"
phone.type =
  addressbook_pb2.Person.HOME
```

Serialization Formats

- [Avro](#)

- Richer data structures
- JSON-specified schema

```
{
  "namespace": "example.avro",
  "type": "record",
  "name": "User",
  "fields": [
    {"name": "name", "type": "string"},
    {"name": "favorite_number", "type": ["int", "null"]},
    {"name": "favorite_color", "type": ["string", "null"]}
  ]
}
```

- [Thrift](#)

- Developed by Facebook
- Now Apache project
- More languages supported
- Supports exceptions and sets

Class Project: TODOs

- Clone / look [DATA605 - Class project](#)
- Look at [Sorrentum GitHub repo](#)
 - Two project examples
 - Price data from Binance
 - Reddit data
- Complete the student [survey](#)
- Take a look at list of projects
- By next lesson:
 - Dr Saggese + TA create teams and assign a project