

# Class Announcements

## 1) Project teams

- Posted [UMD DATA605 - Class Project Teams - Spring 2023](#)
- If your name is not on any of the teams, please send me an email
- No midterm or final exam, complete class project to get a grade

## 2) Team composition

- Teams based on your self-assessed skills
- In each group there should be someone with experience with Git, Docker, Python, and so on
- Teams are not perfect, but none of your team at your future jobs will be
- Working in a team is a skill that takes a long time to hone: let's start practicing it

## 3) Class project complexity

- Projects have about the same complexity
- If not, we will try to account for this when grading
- Projects were assigned randomly to the teams

# Class Announcements

## 4) Add personal info

- Fill the spreadsheet [UMD DATA605 - Class Project Teams - Spring 2023](#) with your information
- Email
- GitHub username (free account at <https://github.com>)
- Telegram Handle for IM (free account at <https://telegram.org>)

## 5) Next steps

- Read carefully [UMD DATA605 - Class Project](#)
- If things are not clear send us an email or add a comment to the Google Doc
- Read the Google doc corresponding to your assigned project from [UMD DATA605 - Class Project Teams - Spring 2023](#)
- Read the first Deliverable 1 from [UMD DATA605 - Class Project](#)

## 6) Golden rule


- *Treat others how you want to be treated*
  - Everybody comes from a different place and different skill level, somebody has a job, somebody has a full-time work
- *If you want to go fast, go alone; if you want to go far, go together*

# **UMD DATA605 - Big Data Systems**

## **DevOps with Docker**

# Docker - Resources

- Concepts in the slides
- Tutorials:
  - [tutorial\\_docker](#)
  - [tutorial\\_docker\\_compose](#)
- We will use Docker during the project
- Web resources
  - [A Beginner-Friendly Introduction to Containers, VMs and Docker](#)
  - [Official Docker Getting Started Tutorial](#)
- Mastery:
  - Poulton, [Docker Deep Dive: Zero to Docker in a single book](#), 2020




# Application Deployment

- For (almost all) Internet companies, **the application is the business**
  - If the application breaks, the business stops working
  - E.g., Amazon, Google, Facebook, on-line banks, travel sites (e.g., Expedia)
- **Problem**
  - How to release / deploy / manage / monitor applications?
- **Solutions**
  - Before 2000s: “bare-metal era”
  - 2000s-2010s: “virtual machine era”
  - > ~2013: “container era”

# DevOps

- **DevOps** = set of practices that combines:
  - software development (dev)
  - IT operations (ops)
- **Containers revolutionized DevOps**
  - Enable true independence between applications and IT ops
    - One team creates an application
    - Another team deploys and manages the applications
  - Create a model for better collaboration (fewer conflicts) and innovation
    - IT: “It doesn’t work!”
    - Devs: “What? It works for me”



# Before Virtualization

- **< 2000s: “Run on bare metal”**
  - Running one or few applications on each server without virtualization
- **Pros**
  - No virtualization overhead
- **Cons**
  - Not safe / not secure since no separation between applications
  - Expensive
- **Expensive / low efficiency**
  - IT would buy a new server for each application
  - Difficult to spec out the machine → buy “big and fast servers”
    - Overpowered servers operating at 5-10% of capacity
    - Tons of money in the 2000 DotCom boom was spent on machines (Sun Servers) and networks (Cisco)



# Virtual Machine Era

- **Circa 2000-2010: Virtual Machine**
  - Virtual machine technology = run multiple copies of OSes on the same hardware
- **Pros**
  - VM runs safely and securely multiple applications on a single server
  - IT could run apps on existing servers with spare capacity
- **Cons**
  - Every VM requires an OS (waste of CPU, RAM, and disk)
  - Monitor and patch each OS
  - Buy an OS license
  - VMs are slow to boot





# Containers Era

- **Circa 2013: Docker becomes ubiquitous**

- **Docker**


- Didn't invent containers
  - Made containers simple and mainstream
- Linux supported containers for some time
    - Kernel namespaces
    - Control groups
    - Union filesystems

- **Pros**

- Containers are fast and portable
- Containers don't require full-blown OS
- All containers run on a single host
- Reduce OS licencing cost
- Reduce overhead of OS patching and maintenance

- **Cons**

- CPU overhead
- Toolchain to learn / use



# Serverless Computing

- **Containers run in an OS, OS runs on a host**

- **Where is the host?**

- Local (your laptop)
    - On premise (your own computer in a rack)
    - Cloud instance (e.g., AWS EC2 instance)

- **What is the host running?**

- Bare-metal server
    - On a virtual machine
    - On a virtual machine running a virtual machine

- **Serverless computing**

- As long your application runs somewhere, you don't care “how” or “where”
  - E.g., AWS Lambda

# HW vs OS Virtualization

- **Hypervisor performs HW virtualization**

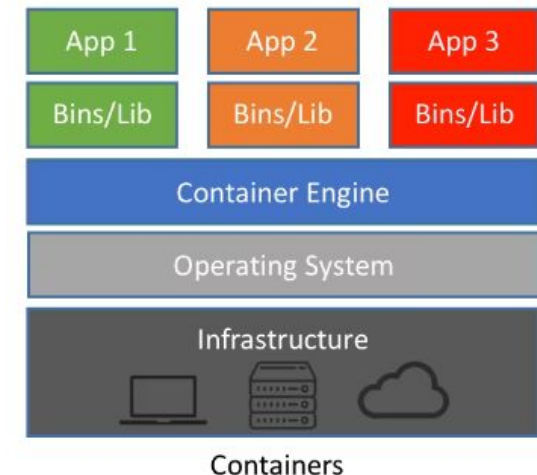
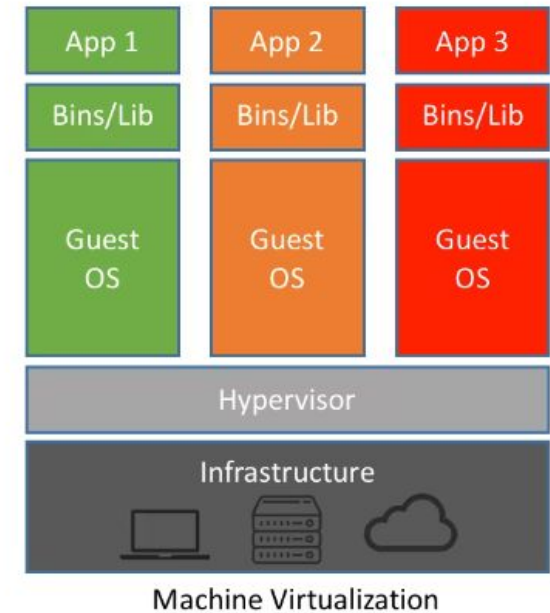
- Carves out physical hardware resources into VMs
- Resources (CPUs, RAM, storage) are allocated to a VM
- It's like having multiple computers

- “Virtual machine tax”

- To run 3 apps, you need 3 VMs
  - Each VM requires time to start
  - Consume CPU, RAM, storage
  - VM needs a OS license
  - VM / OS need admins, patching
- You just want to run 3 apps!

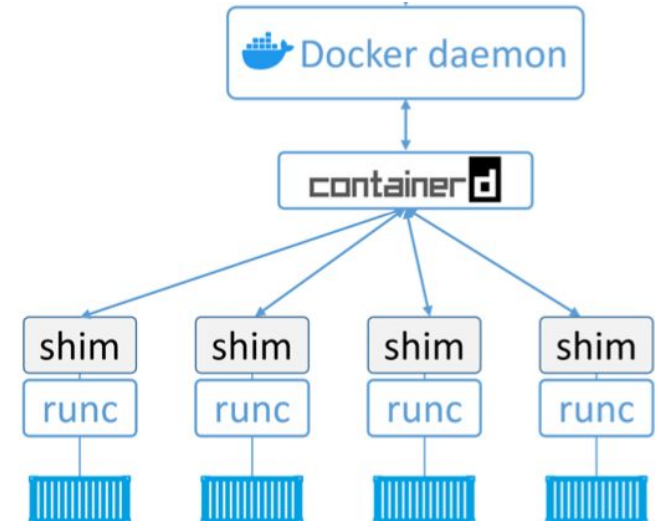

- **Containers perform OS virtualization**

- It's like having multiple OSes



# Docker: Server-Client

- Server-client architecture
- **Docker client**
  - Command line interface
  - Communicate with the server through IPC socket (e.g., `/var/run/docker.sock`) or IP port
- **Docker engine**
  - Run and manage containers
  - Modular and built from several OCI-compliant sub-systems
    - E.g., Docker daemon, **containerd**, **runc**, plug-ins for networking and storage



# Docker Architecture

- **Docker run-time**

- **runc**: start and stop containers
- **containerd**
  - Pull images
  - Create volumes, network interfaces

- **Docker engine**


- **dockerd**
  - Expose remote API
  - Manage images, volumes, networks

- **Docker orchestration**

- **docker swarm**
- Manage clusters of nodes
- Replaced by Kubernetes

- **Open Container Initiative (OCI)**

- Standardize low-level components of container infrastructure
- E.g., image format, run-time API
- “Death” of Docker



# Docker Container

- **Docker Container**

- Lightweight, stand-alone, executable software package
- Include everything needed to run
  - E.g., code, runtime / system libraries, settings
- It is a run-time object
  - vs Docker images are built-time objects
  - Like program running (container) vs program code (image)




# Docker Image

- **Docker Image**

- Unit of deployment
- Contain everything needed by an app to run
  - Application code
  - Application dependencies
  - Minimal OS support
  - File system
- Users can
  - Build images from Dockerfiles
  - Pull pre-built images from a registry
- Multiple layers stacked on top of each other
  - Typically few 100s MBs

# Docker Image Layers

- **Docker image** is a configuration file that lists the layers and some metadata
  - It is composed of read-only layers
  - Each layer is independent from each other
  - Each layer comprising of many files
- **Docker driver**
  - Stacks these layers as a unified filesystem
  - Files from the top layers can obscure the files from the bottom layers
  - Implements a copy-on-write behavior
- **Layer hash**
  - Each layer has an hash based on its content
  - Layers are pulled and pushed compressed
    - A "distribution hash" is used
- **Image hash**
  - Each image has an hash
  - The hash is function of the config file and of the layers
  - When an image changes, a new hash is generated





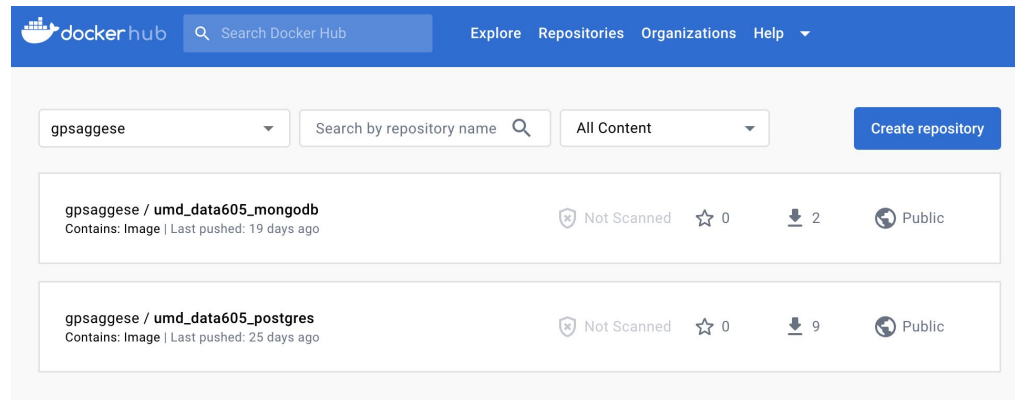
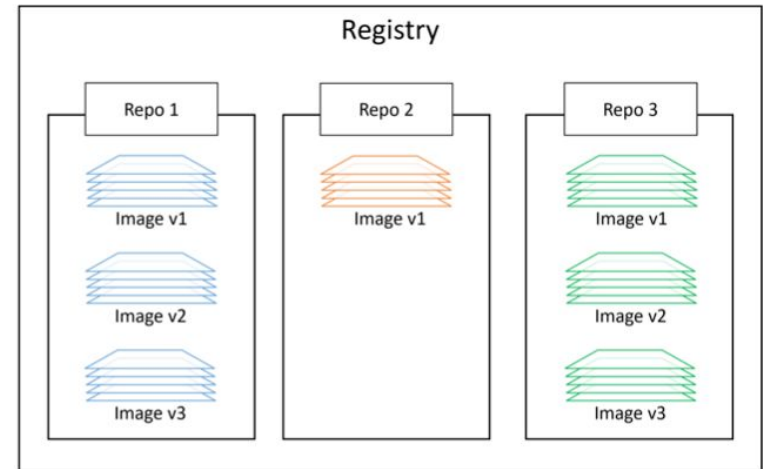
# Docker: Container Data

- A container has access to different data
- **Container storage**
  - It is a copy-on-write layer in the image
  - It is ephemeral (only tmp data)
  - Data inside of containers is persisted as long as the container is not killed
    - If you stop or pause a container data is not lost
  - Containers are designed to be immutable
    - It's not good practice to write “persistent” data into containers
- **Bind-mount a local dir**
  - = a local dir is mounted to a dir inside a container
- **Docker volumes**
  - Docker provides volumes that exist separately from the container
    - E.g., to store the content of a Postgres DB
  - State is permanent across container invocations
  - Can be shared across containers


# Docker Repos

- Docker Repo

- Aka Docker Registry
- Store Docker images
  - `<registry>/<repo>:<tag>`
  - `alpine:latest`
  - E.g., DockerHub, AWS ECR
- Some repos are vetted by Docker
- Unofficial repos shouldn't be trusted
- E.g., <https://hub.docker.com/>



# Devops = Devs + Ops



- **Devs**

- Implement the app
  - E.g., Python, virtual env
- Containerize the app
  - Create Dockerfile
  - Contain the instruction on how to build an image
- Build image
- Run the app as a container
- Test “locally”

- **Ops**

- Download container images
  - Contain filesystem, application, app dependencies
- Start containers
- Destroy containers
- In case of issues, it's easy to reproto the problem
  - “Here is the log”
  - Run command line
  - Deploy on a test system and debug

# Containerizing an App

- = create a container with your app inside
- Develop your application code using the needed dependencies
  - E.g., install dependencies
    - Directly inside a container
    - Inside a virtual env
- Create a Dockerfile describing:
  - your app
  - its dependencies
  - how to run it
- Build image with **docker image build**
- (Optional) Push image to a Docker image registry
- Run / test container from image
- Distribute your app as a container (no installation)

# Building a Container

- **Dockerfile**

- Describe how to create a container

- **Build context**

- `> docker build -t web:latest .`
  - Sent to Docker engine to build the application
  - Directory containing the application and its dependencies
  - Typically the Dockerfile is in the root directory of the build context

# Dockerfile Example

```
FROM python:3.8-slim-buster
```

```
LABEL maintainer="gsaggese@umd.edu"
```

```
WORKDIR /app
```

```
COPY requirements.txt requirements.txt
```

```
RUN pip3 install -r requirements.txt
```

```
COPY . .
```

```
CMD ["python3", "-m", "flask", "run", "--host=0.0.0.0"]
```

# Docker Tutorial

- [tutorial\\_docker.md](#)

# Docker Compose

- **Manage multi-container apps running on a single node**
  - Describe the app in a single declarative configuration YAML file
    - Instead of scripts with long Docker commands
  - Compose talks to Docker API to achieve what you requested
  - E.g., you need a client app and Postgres DB
  - E.g., microservices
    - Web front-end
    - Ordering
    - Back-end DB
- In 2020 Docker Compose has become an open standard for “code-to-cloud” process
- **Manage multi-container apps running on multiple hosts**
  - Docker Stacks / Swarm
  - Kubernetes



# Docker Compose: Commands

```
> docker compose --help
```

Usage: `docker compose [OPTIONS] COMMAND`

Options:

<code>--env-file string</code>	Specify an alternate environment file.
<code>-f, --file stringArray</code>	Compose configuration files
<code>-p, --project-name string</code>	Project name

Commands:

<code>build</code>	Build or rebuild services
<code>convert</code>	Converts the compose file to platform's canonical format
<code>cp</code>	Copy files/folders between a service container and the local filesystem
<code>create</code>	Creates containers for a service.
<code>down</code>	Stop and remove containers, networks
<code>events</code>	Receive real time events from containers.
<code>exec</code>	Execute a command in a running container.
<code>images</code>	List images used by the created containers
<code>kill</code>	Force stop service containers.
<code>logs</code>	View output from containers
<code>ls</code>	List running compose projects
<code>pause</code>	Pause services
<code>port</code>	Print the public port for a port binding.
<code>ps</code>	List containers
<code>pull</code>	Pull service images
<code>push</code>	Push service images
<code>restart</code>	Restart containers
<code>rm</code>	Removes stopped service containers
<code>run</code>	Run a one-off command on a service.
<code>start</code>	Start services
<code>stop</code>	Stop services
<code>top</code>	Display the running processes
<code>unpause</code>	Unpause services
<code>up</code>	Create and start containers
<code>version</code>	Show the Docker Compose version information

# Docker Compose: Tutorial Example

- The default name for a Compose file is ``docker-compose.yml``
  - You can specify ``-f`` for custom filenames
- **Top-level keys** are:
  - ``version``:
    - Mandatory first line to specify API version
    - Ideally always use the latest version
    - Typically 3 or higher
  - ``services``:
    - Define the different microservices
  - ``networks``:
    - Creates new networks
    - By default it creates a ``bridge`` network to connect multiple containers on the same Docker host
  - ``volumes``:
    - Creates new volumes
- **Key in services** describe a different “service” in terms of container
  - **Inner keys** specify the params of Docker run command

**version:** "3.8"

**services:**

**web-fe:**

**build:** .

**command:** python app.py

**ports:**

– target: 5000

published: 5000

**networks:**

– counter-net

**volumes:**

– type: volume

source: counter-vol

target: /code

**redis:**

image: "redis:alpine"

networks:

counter-net:

**networks:**

counter-net:

**volumes:**

counter-vol:

# Docker Compose: Tutorial

- Example taken from <https://github.com/nigelpoulton/counter-app>
- [tutorial\\_docker\\_compose](#)
  - > cd tutorials/tutorial\_docker\_compose
  - > vi tutorial\_docker\_compose.md