

UMD DATA605 - Big Data Systems

NoSQL Document Stores

MongoDB

CouchDB

Dr. GP Saggese

gsaggese@umd.edu

with thanks to Profs

Alan Sussman (UMD)

Amol Deshpande (UMD)

Nguyen Vo (Utah State U.)

Kathleen Durant (Northeastern U.)

UMD DATA605 - Big Data Systems

NoSQL Document Stores

MongoDB

Couchbase

Key-Value Store vs Document DBs

- **Key-value stores**

- Basically a map or dictionary
 - E.g., HBase, Redis
- Typically only look up values by key
 - Sometimes can do search in value field with a pattern
- Uninterpreted value (e.g., binary blob) associated with a key
- Typically one namespace for all key-values

- **Document DBs**

- Collect sets of key-value pairs into *documents*
 - E.g., MongoDB, CouchDB
- Documents typically represented in JSON, XML, or binary JSON
- Documents organized into *collections*, similar to tables in relational DBs
 - Large collections can be partitioned and indexed

UMD DATA605 - Big Data Systems


NoSQL Document Stores

MongoDB

Couchbase

Resources

- Concepts in slides
- MongoDB tutorial
- Web
 - <https://www.mongodb.com/>
 - [Official docs](#)
 - [pymongo](#)
- Book
 - [Seven Databases in Seven Weeks, 2e](#)



MongoDB



- Developed by MongoDB Inc.
 - Founded in 2007
 - Based on DoubleClick experience with large-scale data
 - Mongo comes from “hu-mongo-us”
- One of the most used NoSQL DBs (if not the most used)
- Document-oriented NoSQL database
 - Schema-less
 - No Data Definition Language (DDL)
 - In practice, you can store maps with any keys and values that you choose
 - Application tracks the schema, mapping between documents and their meaning
 - Keys are hashes stored as strings
 - Document Identifiers `_id` created for each document (field name reserved by system)
 - Values use BSON format
 - Based on JSON (B stands for Binary)
- Written in C++
- Supports APIs (drivers) in many languages
 - E.g., JavaScript, Python, Ruby, Java, Scala, C++, ...

MongoDB: Example of Document

- A **document** is a JSON data structure
- Correspond to a row in a relational DB
 - without schema
 - values nested to an arbitrary depth
 - primary key is `_id`

```
{
  "_id" : ObjectId("4d0b6da3bb30773266f39fea"),
  "country" : {
    "$ref" : "countries",
    "$id" : ObjectId("4d0e6074deb8995216a8309e")
  },
  "famous_for" : [
    "beer",
    "food"
  ],
  "last_census" : "Sun Jan 07 2018 00:00:00 GMT -0700 (PDT)",
  "mayor" : {
    "name" : "Ted Wheeler",
    "party" : "D"
  },
  "name" : "Portland",
  "population" : 582000,
  "state" : "OR"
}
```

MongoDB: Functionalities

- Design goals:
 - Performance
 - Scalability
 - Rich data access
- Dynamic schema
 - No DDL (Data Definition Language)
 - Secondary indexes
- Query language via an API
- Several levels of consistency
 - E.g., atomic writes and fully-consistent reads
- No joins nor transactions across multiple documents
 - Makes distributed queries easy and fast
- High availability through replica sets
 - E.g., primary replication with automated failover
- Built-in sharding
 - Horizontal scaling via automated range-based partitioning of data
 - Reads and writes distributed over shards

Relational DBs vs MongoDB: Terms and Concepts

RDBMS Concepts	MongoDB Concepts	Meaning in MongoDB
database	database	Container for collections
relation / table, view	collection	Group of documents
row / instance	document (BSON)	Group of fields
column / attributes	field	A name-value pair
index	index	
primary keys	<code>_id</code> field	Always the primary key
foreign key	reference	
table joins	embedded documents and linking	

Relational vs Document DB Workflows

- **Relational DBs**

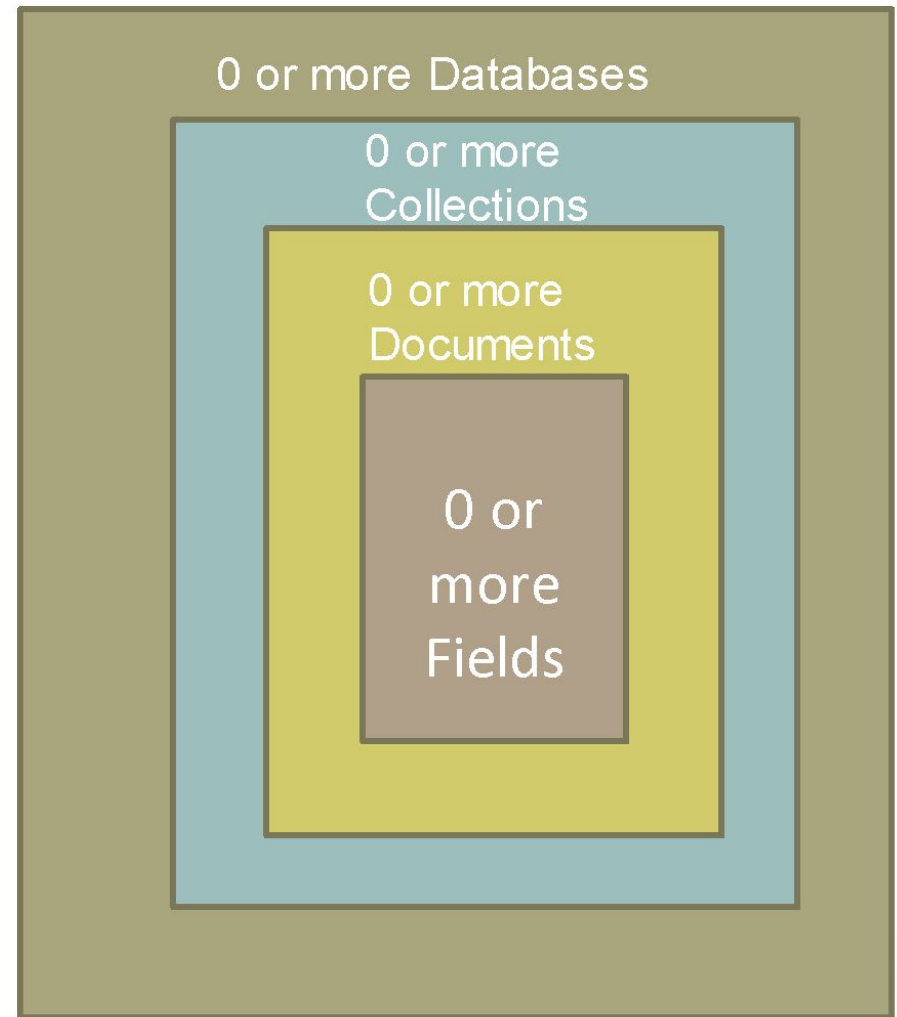
- E.g., PostgreSQL
- Know what you want to store
 - Tabular data
- Do not know how to use it
 - Static schema allows query flexibility (e.g., joins)
- Complexity is at insertion time
 - Decide how to represent the data

- **Document DBs**

- E.g., MongoDB
- No assumptions on what to store
 - E.g., irregular JSON data
- Know a bit how to access
 - E.g., it's a nested key-value map
- Complexity is at access time
 - Get the data from the server
 - Process data on the client side


MongoDB: Hierarchical Objects

- An **instance** has:
 - zero or more “databases”
 - same as Postgres
- A **database** has:
 - zero or more “collections”
 - ~ Postgres tables
- A **collection** has:
 - zero or more “documents”
 - ~ Postgres rows
- A **document** has:
 - one or more “fields”
 - It has always the `_id`
 - ~ Postgres columns



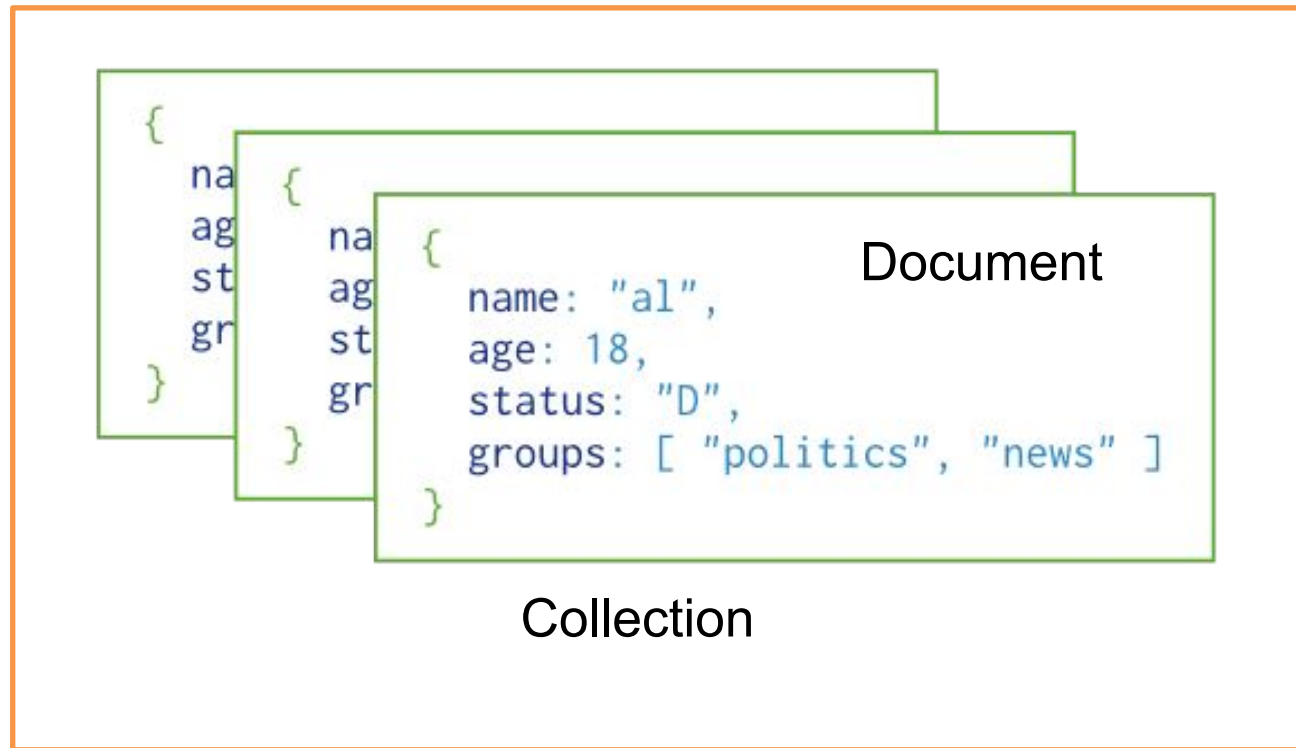
Why Use MongoDB?

- Simple to query
 - Do the work on client side
- It's fast
 - 2-10x faster than Postgres
- Data model / functionalities are applicable to most web applications
 - Semi-structured data
 - Quickly evolving systems
- Easy and fast integration of data
- Not well suited for heavy and complex transactions systems



MongoDB Data Model

A **collection** (~table) includes **documents** (~rows)



From <https://www.mongodb.com/docs/manual/core/data-modeling-introduction>

MongoDB Data Model

- **Documents** are composed of field and value pairs
 - Field names are strings
 - Each value is any BSON type (native data types, other documents, arrays of documents)
- E.g.,
 - `_id` holds an ObjectId
 - `name` holds a document that contains the fields `first` and `last`
 - `birth` and `death` are of Date type
 - `contribs` holds an array of strings
 - `views` holds a value of the NumberLong type

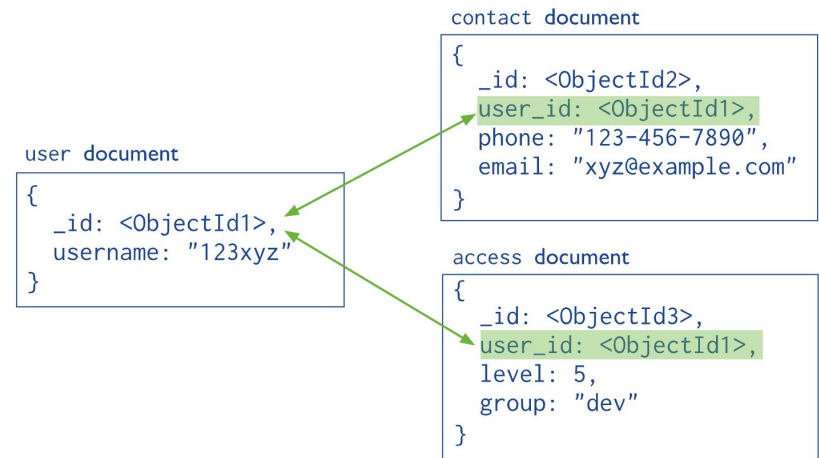
```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```

← field: value
← field: value
← field: value
← field: value

```
{  
  _id: ObjectId("5099803df3f4948bd2f98391"),  
  name: { first: "Alan", last: "Turing" },  
  birth: new Date('Jun 23, 1912'),  
  death: new Date('Jun 07, 1954'),  
  contribs: [ "Turing machine", "Turing test", "Turingery" ],  
  views : NumberLong(1250000)  
}
```

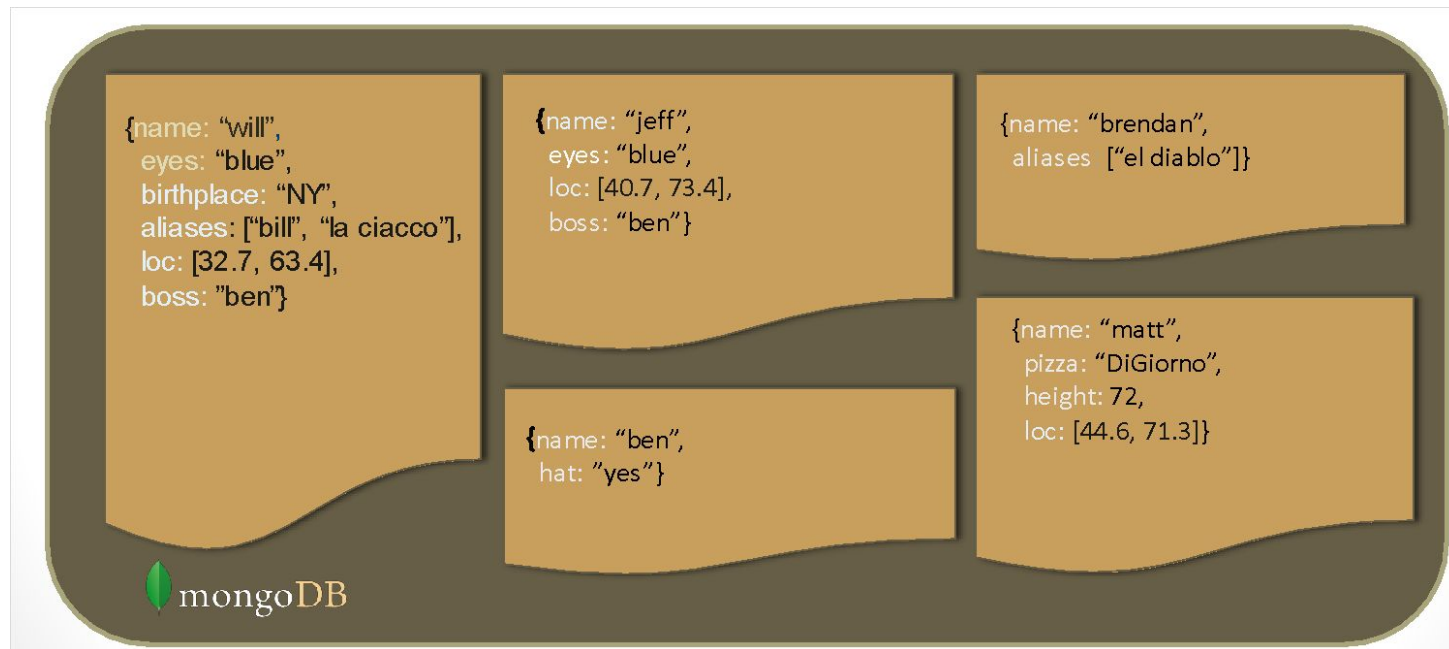
MongoDB Data Model

- Documents can be nested
- **Denormalized data models**
 - Allow to retrieve and manipulate data in a single operation
 - Store multiple related pieces of information in the same record
- **Normalized data models**
 - Eliminate duplication
 - Represent many-to-many relationships



Schema Free

- MongoDB does not need any pre-defined data schema
- Every document in a collection can have different fields and values
 - No need for NULL data fields



JSON Format


- JavaScript Object Notation
- Data is in name / value pairs
- A name / value pair consists of a field name followed by a colon, followed by a value
`"name": "R2-D2"`
- Data is separated by commas
`"name": "R2-D2", race : "Droid"`
- Curly braces {} hold objects
`{"name": "R2-D2", race : "Droid", affiliation: "rebels"}`
- An array is stored in brackets []
`[{"name": "R2-D2", race: "Droid", affiliation: "rebels"}, {"name": "Yoda", affiliation: "rebels"}]`
- Support embedding of nested objects within other objects, or just references

BSON Format

- Binary-encoded serialization of JSON-like documents
 - <https://bsonspec.org>
- Zero or more key/value pairs are stored as a single entity
- Each entry consists of:
 - a field name
 - a data type
 - a value
- Large elements in a BSON document are prefixed with a length field to facilitate scanning
- MongoDB understands the internals of BSON objects, even nested ones
 - Can build indexes and match objects against query expressions for BSON keys

ObjectId

- Each JSON data contains an `_id` field of type ObjectId
 - Like a `SERIAL` incrementing a numeric primary key in PostgreSQL
- An ObjectId is always 12 bytes, composed of:
 - a timestamp
 - client machine ID
 - client process ID
 - a 3-byte auto-incremented counter
- Each Mongo process can handle its own ID generation without colliding
 - Mongo has a distributed nature
- Details [here](#)



MongoDB Features

- Document-oriented NoSQL store
- Rich querying
 - Full index support
- Fast in-place updates
- Agile and scalable
 - Replication and high availability
 - Auto-Sharding
 - Map-reduce functionality
- Scale horizontally over commodity hardware
 - Lots of relatively inexpensive servers

MongoDB vs Relational DBs

- Keep the functionality that works well in RDBMSs
 - Ad-hoc queries
 - Fully featured indexes
 - Secondary indexes
- Do not offer RDBMS functionalities that don't distribute
 - Long running multi-row transactions
 - Joins
 - Both artifacts of the relational data model (row x column)

Indexes

- Primary index automatically created on the `_id` field
 - B+ tree indexes
- Users can create secondary indexes to:
 - Improve query performance
 - Enforce unique values for a particular field
- Support single field index and compound index
 - Order of the fields in a compound index matters (like SQL)
 - If you index a field that holds an array value, MongoDB creates separate index entries for every element of the array
- Sparse property of an index
 - The index contains only entries for documents that have the indexed field
 - Ignore records that do not have the field defined
- If an index is unique and sparse
 - Reject records that have a duplicate key value
 - Allow records that do not have the indexed field defined
- Details at <https://www.mongodb.com/docs/manual/indexes/>

CRUD Operations

- CRUD = Create, Read, Update, Delete
- **Create**
 - `db.collection.insert(<document>)`
 - `db.collection.save(<document>)`
 - `db.collection.update(<query>, <update>, {upsert: true})`
- **Read**
 - `db.collection.find(<query>, <projection>)`
 - `db.collection.findOne(<query>, <projection>)`
- **Update**
 - `db.collection.update(<query>, <update>, <options>)`
- **Delete**
 - `db.collection.remove(<query>, <justOne>)`

Details at <https://www.mongodb.com/docs/manual/crud/>

Create Operations

- `db.collection` specifies the collection (like a 'table') to store the document

`db.collection_name.insert(<document>)`

- Omit the `_id` field to have MongoDB generate a unique key

`db.parts.insert({type: "screwdriver", quantity: 15})`

`db.parts.insert({_id: 10, type: "hammer", quantity: 1})`

- Update 1 or more records in a collection satisfying *query*

`db.collection_name.update(<query>, <update>, {upsert: true})`

- Update an existing record or creates a new record

`db.collection_name.save(<document>)`

Read Operations

- find provides functionality similar to the SQL SELECT command, with
 - <query> = where condition
 - <projection> = fields in result set**db.collection.find(<query>, <projection>).cursor**
- **db.parts.find({parts: "hammer"}).limit(5)**
 - Has cursors to handle a result set
 - Can modify the query to impose limits, skips, and sort orders
 - Can specify to return the 'top' number of records from the result set
- **db.collection.findOne(<query>, <projection>)**

More Query Examples

SQL

```
SELECT * FROM users WHERE age>33
```

```
SELECT * FROM users WHERE age!=33
```

```
SELECT * FROM users WHERE name LIKE  
"%Joe%"
```

```
SELECT * FROM users WHERE a=1 and b='q'
```

```
SELECT * FROM users WHERE a=1 or b=2
```

```
SELECT * FROM foo WHERE name='bob' and  
(a=1 or b=2 )
```

```
SELECT * FROM users WHERE age>33 AND  
age<=40
```

Mongo

```
db.users.find({age:{$gt:33}})
```

```
db.users.find({age:{$ne:33}})
```

```
db.users.find({name:/Joe/})
```

```
db.users.find({a:1,b:'q'})
```

```
db.users.find({$or: [{a: 1}, {b: 2}]})
```

```
db.foo.find({name: "bob", $or :  
[ {a: 1}, {b: 2} ]})
```

```
db.users.find({'age':{$gt:33,$lte:40}})
```

Query Operators

Command	Description
\$regex	Match by any PCRE-compliant regular expression string (or just use the // delimiters as shown earlier)
\$ne	Not equal to
\$lt	Less than
\$lte	Less than or equal to
\$gt	Greater than
\$gte	Greater than or equal to
\$exists	Check for the existence of a field
\$all	Match all elements in an array
\$in	Match any elements in an array
\$nin	Does not match any elements in an array
\$elemMatch	Match all fields in an array of nested documents
\$or	or
\$nor	Not or
\$size	Match array of given size
\$mod	Modulus
\$type	Match if field is a given datatype
\$not	Negate the given operator check

Update Operations

- `db.collection_name.insert(<document>)`
 - Omit the `_id` field to have MongoDB generate a unique key

```
db.parts.insert( {type: "screwdriver", quantity: 15 } )  
db.parts.insert({_id: 10, type: "hammer", quantity: 1 } )
```
- `db.collection_name.save(<document>)`
 - Updates an existing record or creates a new record
- `db.collection_name.update(<query>, <update>, {upsert: true})`
 - Will update 1 or more records in a collection satisfying query
- `db.collection_name.findAndModify(<query>, <sort>, <update>, <new>, <fields>, <upsert>)`
 - Modify existing record(s) – retrieve old or new version of the record

Delete Operations

- `db.collection_name.remove(<query>, <justone>)`
 - Delete all records from a collection or matching a criterion
 - `<justone>` specifies to delete only 1 record matching the criterion
- Remove all parts starting with h
`db.parts.remove(type: /^h/ }`
- Delete all documents in the parts collections
`db.parts.remove()`

MongoDB Tutorial

Tutorial is at [GitHub](#)

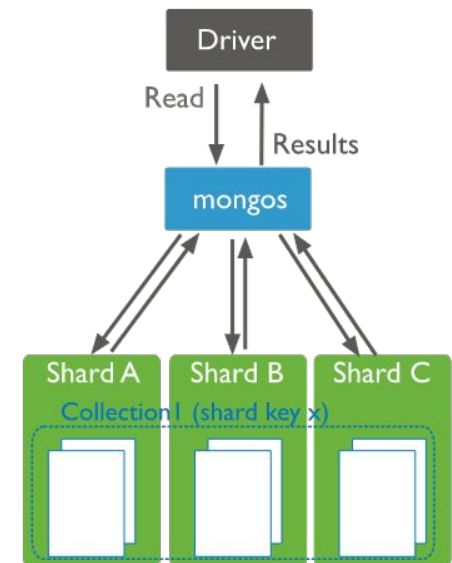
The instructions are [here](#)

```
> cd $GIT_REPO/tutorials/tutorial_mongodb
```

```
> vi tutorial_mongo.md
```

MongoDB Processes and Configuration

- ***mongod***: database instance (i.e., the server process)
- ***mongosh***: an interactive shell (i.e., a client)
 - Fully functional JavaScript environment for use with a MongoDB
- ***mongos***: database router
 - Process all requests
 - Decide how many and which *mongods* should receive the query (sharding / partitioning)
 - Collate the results
 - Send result back to the client
- You should have:
 - One *mongos* for the whole system no matter how many *mongods* you have; or
 - One local *mongos* for every client if you wanted to minimize network latency



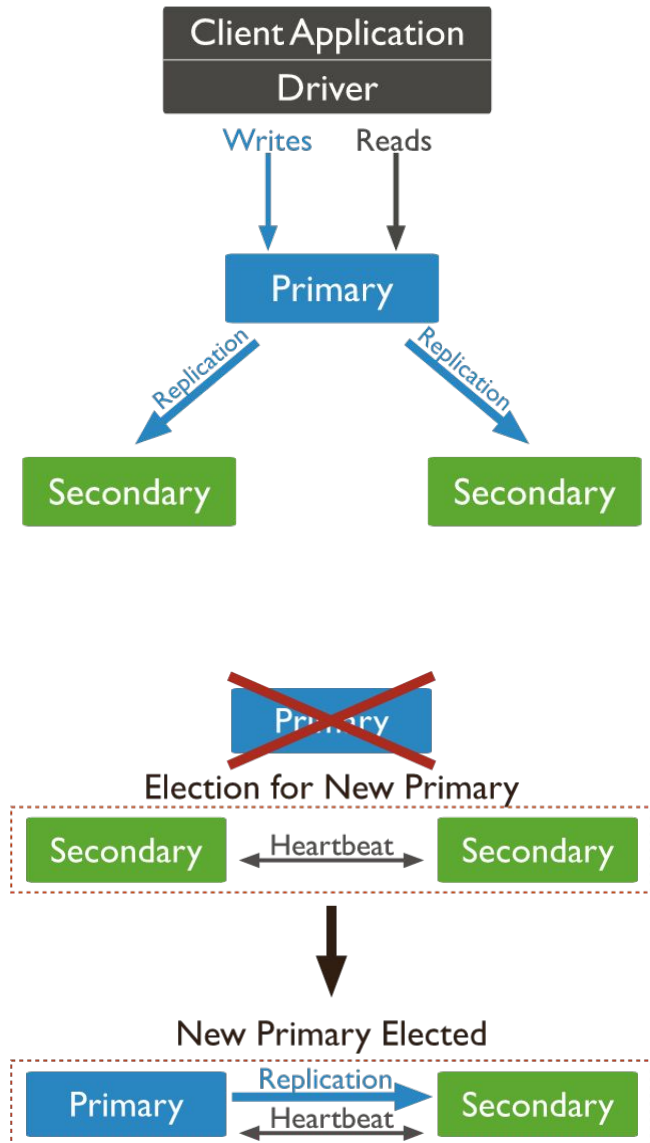
MapReduce Functionality

- Perform aggregator functions given a collection of (keys, value) pairs
- Must provide at least a map function, reduction function, and the name of the result set

```
db.collection.mapReduce(  
  <mapfunction>,  
  <reducefunction>,  
  {  
    out: <collection>,  
    query: <document>,  
    sort: <document>,  
    limit: <number>,  
    finalize: <function>,  
    scope: <document>,  
    jsMode: <boolean>,  
    verbose: <boolean>  
  })
```


Data Replication

- Ensure redundancy, backup, and automatic failover
- Replication occurs through groups of servers known as **replica sets** (for each shard)
 - **Primary set:** set of servers that client asks direct updates to
 - **Secondary set:** set of servers used for duplication of data
 - Different properties can be associated with a secondary set, e.g., secondary-only, hidden delayed, arbiters, non-voting
- If the primary set fails the secondary sets “vote” to elect the new primary set

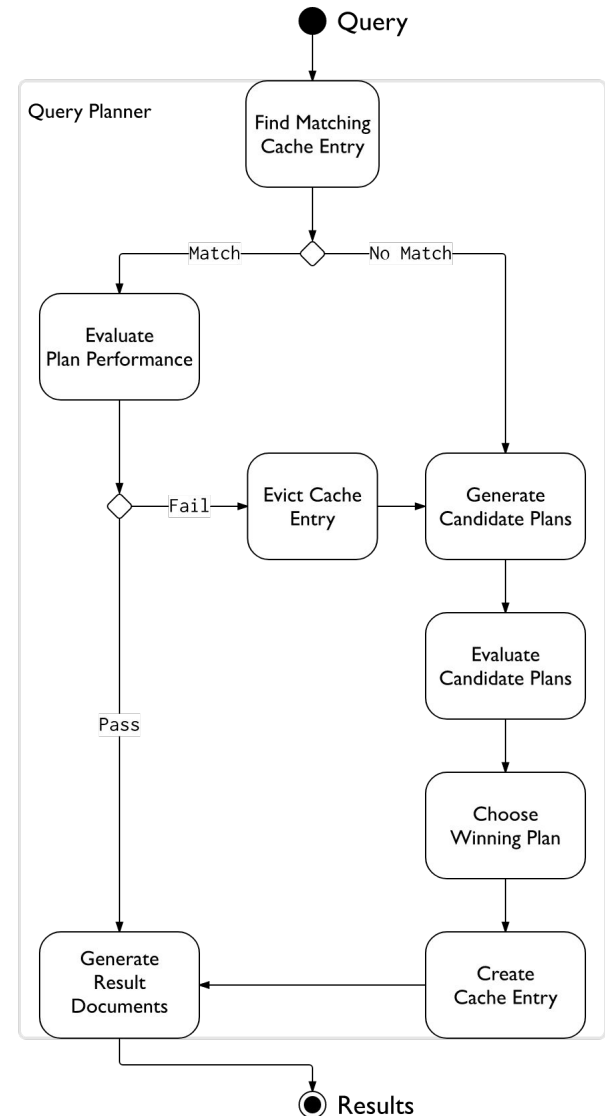


Data Consistency

- Client decides how to enforce consistency for reads
 - All writes and *consistent* reads go to the primary
 - All *eventually consistent* reads are distributed among the secondaries
- Reads to a primary have strict consistency
 - Reads reflect the latest changes to the data
- Reads to a secondary have eventual consistency
 - Updates propagate gradually
 - Client may read a previous state of the database

Query Optimizer

- Not static and cost-based optimizer (like in RDBMSs)
- MongoDB optimizer tries different query plans and learns which ones perform well
 - Because there are no joins, the space of query plans is not so large
 - When testing new plans
 - Execute multiple query plans in parallel
 - As soon as one plan finishes, terminate the other plans
 - If a plan that was working well starts performing poorly try again different plans
 - E.g, data in the DB has changed, parameter values to a query are different

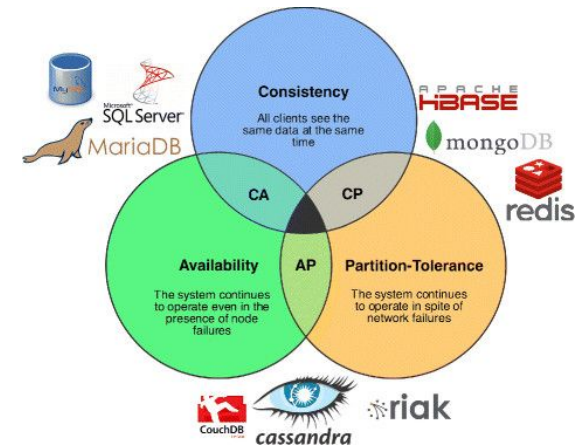


Strengths

- Provide a query language
- High-performance
 - Implemented in C++
- Very rapid development, open source
 - Support for many platforms
 - Many language drivers
- Built to address a distributed database system
 - Sharding
 - Replica sets of data
- Tunable consistency
- Useful when working with a huge quantity of data not requiring a relational model
 - What really matters is the ability to store and retrieve great quantities of data
 - The relationships between the elements does not matter

Limitations

- No referential integrity
- Lack transactions, joins
- High degree of denormalization
 - Updating something in many places instead of one
- Lack of predefined schema is a double-edged sword
 - You must have a data model in your application
 - Objects within a collection can be completely inconsistent in their fields
- CAP Theorem: targets consistency and partition tolerance



UMD DATA605 - Big Data Systems

NoSQL document stores

MongoDB

Couchbase

Couchbase

- NoSQL document-oriented DB like MongoDB
- Couchbase = merge of CouchDB and membase
 - *CouchDB*
 - Open source document store
 - HTTP RESTful API for reading and updating (add, edit, delete) documents
 - Support all 4 ACID properties
 - *membase*
 - Distributed key-value store
 - Designed to scale both up and down
 - Highly available and partition tolerant
- Couchbase
 - Uses HTTP protocol to query and interact with objects in the DB
 - Objects stored in *buckets*
 - Just a collection of documents (in JSON), with no special relation to one another
- For CAP, get consistency and partition tolerance by default, high availability instead of consistency through use of multiple clusters



Architecture

- Every Couchbase node consists of:
 - a data service
 - index service
 - query service
 - cluster manager component
- Services can run on separate nodes of the cluster, if needed
- Data replication across nodes of a cluster (and across data centers)
- Data manager *asynchronously* writes data to disk after acknowledging to the client
 - Optionally can ensure data is written to more than one server before acknowledging a write

Queries

- Can create multiple views over documents
 - Views are optimized/indexed by Couchbase for fast queries
 - Only re-indexed when underlying documents change a lot
- Perform well when there are infrequent changes to the structure of documents
 - And know in advance what kinds of queries you want to execute
- Uses a custom query language called N1QL, based on SQL
 - But runs on JSON documents
 - And queries over multiple documents using joins
- Can also do full-text searches using the indexes
- Map-reduce support
 - First define a view with the columns of the document your are interested in
 - Called the *map*
 - Optionally define aggregate functions over the data
 - The *reduce* step

Couchbase vs MongoDB

- According to Couchbase advocates
- **MongoDB**: hard to scale from single replica set to fully distributed environment
- **MongoDB**: performance degrades with increasing numbers of clients/users
 - **Couchbase**: Scales seamlessly, with an in-memory architecture, and able to scale across multiple nodes
- **MongoDB**: susceptible to data loss from failures
 - **Couchbase**: no master, no single point of failure
 - During failover, prevents different nodes from accepting simultaneous reads of writes of same data (maintains consistency)
- **MongoDB**: require a 3rd party cache to help it perform well
 - **Couchbase**: has integrated in-memory cache (memcached)
 - Keeps frequently accessed documents, metadata, and indexes in RAM, yielding high read/write throughput at low latency