

UMD DATA605 - Big Data Systems


Data Wrangling

Dr. GP Saggese
gsaggese@umd.edu

with thanks to Amol Deshpande

Resources

- [Pandas tutorial](#)
- Class project
- Web
 - <https://pandas.pydata.org>
 - Onslaught of free resources
- Mastery
 - <https://wesmckinney.com/book>



Overview

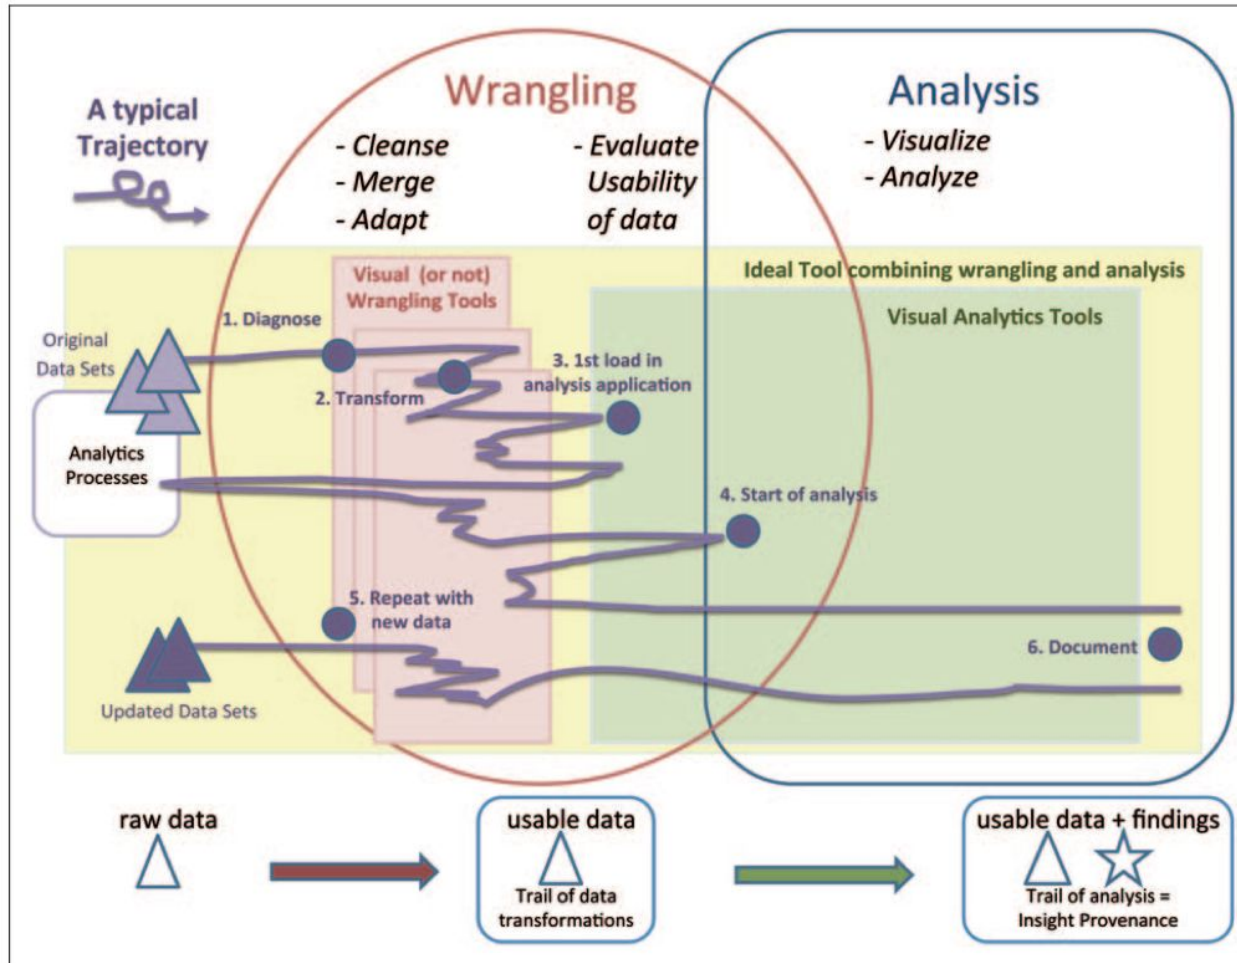
- **Data wrangling**

- Get data into a structured form suitable for analysis
- Aka data preparation, data munging, data curation
- Often it is the step where majority of time (80-90%) is spent

- **Key steps**

- Scraping: extract information from sources (e.g., webpages, spreadsheets)
- Data transformation: get data into the right structure
- Data integration: combine data from multiple sources
- Information extraction: extract structured information from unstructured / text sources
- Data cleaning: remove inconsistencies / errors

Overview



Overview

- Many of the data wrangling problems are not easy to formalize, and have seen little research work, e.g.,
 - Data transformation, i.e., put the data in the “right” structure
 - Information extraction: highly domain specific
 - Data cleaning: somewhat studied (e.g., tidy data)
- Others aspects of integration have been studied in depth, e.g.,
 - Schema mapping
 - Data integration
- Typical workflow
 - From [Data Cleaning: Problems and Current Approaches](#)
 - Somewhat old: data is mostly coming from structured sources
 - For a data scientist, the data scraping is equally important

Overview

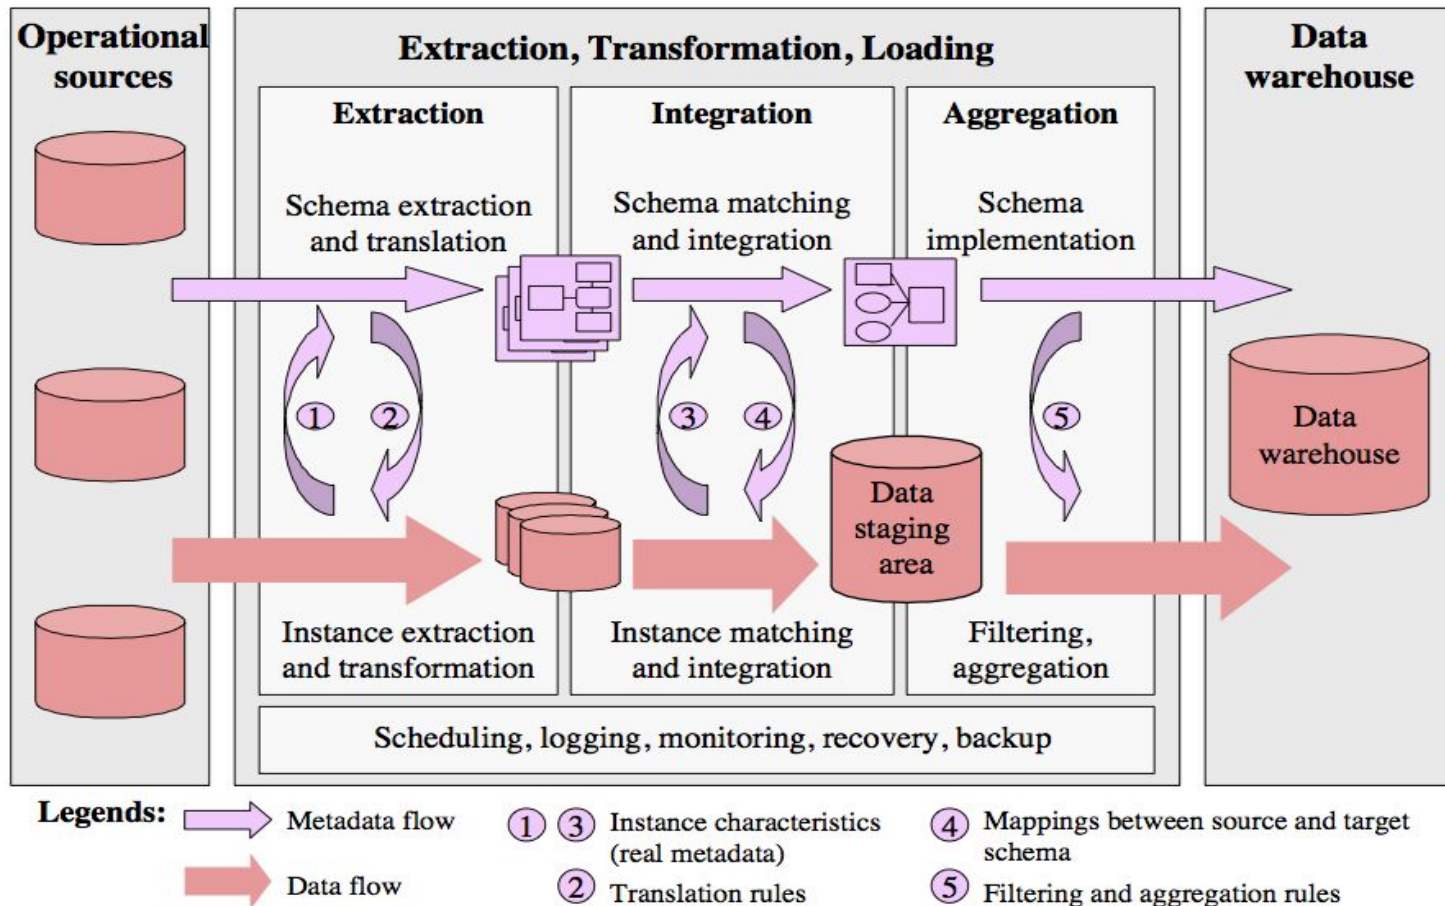


Figure 1. Steps of building a data warehouse: the ETL process

Data Scraping

- Data may reside in a wide variety of different sources
 - Files (e.g., CSV, JSON, XML)
 - Different databases
 - Spreadsheets
- Most analytical tools support importing data from such sources
 - Adapters to load data
- Web scraping: scraping data from web sources is tougher
 - In some cases there may be APIs
 - In other cases data may have to be explicitly scraped
 - Often pipelines are set up to do this on a periodic basis
 - Can be fragile
 - Several tools out there to do this (somewhat) automatically
 - E.g., import.io, [portia](https://portia.io), ...

Tidy Data

- Tidy data, Wickham, 2014 ([here](#))
 - Each variable forms a column
 - Each observation forms a row
- Wide vs long format

	treatmenta	treatmentb
John Smith	—	2
Jane Doe	16	11
Mary Johnson	3	1

	John Smith	Jane Doe	Mary Johnson
treatmenta	—	16	3
treatmentb	2	11	1

“Messy” data

name	trt	result
John Smith	a	—
Jane Doe	a	16
Mary Johnson	a	3
John Smith	b	2
Jane Doe	b	11
Mary Johnson	b	1

Tidy data

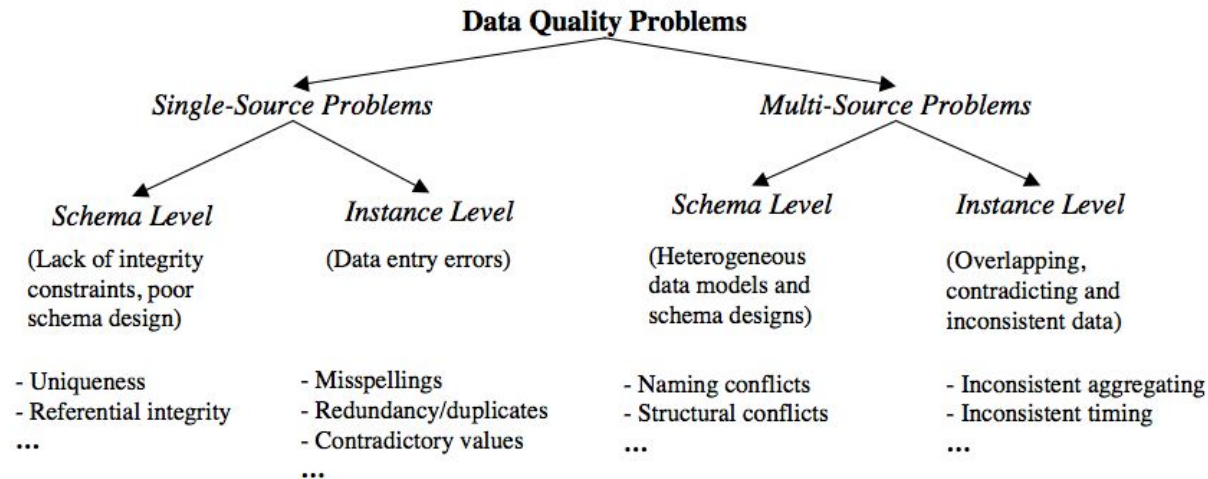
type	date	clicks	conversions	impressions
0	2020-01-01	1.0	NaN	18.0
1	2020-01-02	2.0	NaN	19.0
2	2020-01-03	1.0	1.0	14.0
3	2020-01-04	NaN	NaN	5.0
4	2020-01-05	1.0	NaN	8.0
5	2020-01-06	1.0	1.0	15.0
6	2020-01-07	2.0	NaN	8.0

Wide format

	date	type	count
0	2020-01-01	impressions	18.0
1	2020-01-02	impressions	19.0
2	2020-01-03	impressions	14.0
3	2020-01-04	impressions	5.0
4	2020-01-05	impressions	8.0
...
91	2020-01-28	conversions	NaN
92	2020-01-29	conversions	NaN
93	2020-01-30	conversions	NaN
94	2020-01-31	conversions	NaN
95	2020-02-01	conversions	NaN

Long format

Data Quality Problems



Scope/Problem		Dirty Data	Reasons/Remarks
Attribute	Missing values	phone=9999-999999	unavailable values during data entry (dummy values or null)
	Misspellings	city="Liipzig"	usually typos, phonetic errors
	Cryptic values, Abbreviations	experience="B"; occupation="DB Prog."	
	Embedded values	name="J. Smith 12.02.70 New York"	multiple values entered in one attribute (e.g. in a free-form field)
	Misfielded values	city="Germany"	
Record	Violated attribute dependencies	city="Redmond", zip=77777	city and zip code should correspond
Record type	Word transpositions	name ₁ = "J. Smith", name ₂ ="Miller P."	usually in a free-form field
	Duplicated records	emp ₁ =(name="John Smith",...); emp ₂ =(name="J. Smith",...)	same employee represented twice due to some data entry errors
	Contradicting records	emp ₁ =(name="John Smith", bdate=12.02.70); emp ₂ =(name="John Smith", bdate=12.12.70)	the same real world entity is described by different values
Source	Wrong references	emp=(name="John Smith", deptno=17)	referenced department (17) is defined but wrong

Single-Source Problems

- Depends largely on the source
- Databases can enforce constraints
- Data extracted from files or spreadsheets is often clean
- Data scraped from web-pages is much more messy
- Types of problems:
 - Ill-formatted data (especially from web-pages or files or spreadsheets)
 - Missing or illegal values, misspellings, use of wrong fields, extraction issues (e.g., not easy to separate out different fields)
 - Duplicated records, contradicting information, referential integrity violations
 - Unclear default values
 - Evolving schemas or classification schemes (for categorical attributes)
 - Outliers

Multi-Source Problems

- Different sources are:
 - Developed separately
 - Maintained by different people
 - Stored in different systems
- Issue 1: Schema mapping / transformation
 - Mapping information across sources
 - Naming conflicts: same name used for different objects
 - Structural conflicts: different representations across sources
- Issue 2: Entity resolution
 - Matching entities across sources
- Issue 3: Data quality issues
 - Contradicting information
 - Mismatched information
 - ...

Data Cleaning: Outlier Detection

- Quantitative Data Cleaning for Large Databases, Hellerstein, 2008 ([here](#))
 - Focuses on quantitative data (i.e., integers/floats that measure some quantities of interest)
- Sources of errors in data
 - Data entry errors: users putting in arbitrary values to satisfy the form
 - Measurement errors: especially sensor data
 - Distillation errors: errors that pop up during processing and summarization
 - Data integration errors: inconsistencies across sources that are combined together

Univariate Outlier Detection


- A set of values can be characterized by metrics such as center (e.g., mean), dispersion (e.g., standard deviation), and skew
- Can be used to identify outliers
 - Must watch out for "masking": one extreme outlier may alter the metrics sufficiently to mask other outliers
 - Use robust statistics: considers effect of corrupted data values on distributions
 - Robust center metrics: median, k% trimmed mean (i.e., discard lowest and highest k% values)
 - Robust dispersion: median absolute deviation (MAD), median distance of all the values from the median value
- A reasonable approach to find outliers (assuming normal distribution)
 - Any data points $1.4826 \times \text{MAD}$ away from median
 - May need to eyeball the data (e.g., plot a histogram) to decide if this is true

Outlier Detection

- [Wikipedia Article on Outliers](#) lists several other normality-based tests for outliers
- If data appears to be not normally distributed:
 - Distance-based methods: look for data points that do not have many neighbors
 - Density-based methods:
 - Define *density* to be average distance to k nearest neighbors
 - *Relative density* = density of node/average density of its neighbors
 - Use relative density to decide if a node is an outlier
- Most of these techniques start breaking down as the dimensionality of the data increases
 - *Curse of dimensionality*
 - Can project data into lower-dimensional space and look for outliers there
 - Not as straightforward

Multivariate Outliers

- Analogous to univariate
- One set of techniques based on assuming data follows a *multi-variate normal distribution*
 - Defined by a *mean* μ and a *covariance matrix* Σ
- Mahalanobis distance of a point
 - Square root of $(x - \mu)' \Sigma^{-1} (x - \mu)$
 - Measures how far the point x is from the multivariate normal distribution
 - Outliers are points that are too far away
- Mean / covariance are not *robust* (sensitive to outliers)
 - Iterative approach: remove points with high Mahalanobis distance, recompute the mean and covariance
 - Several other general approaches discussed in the reference by Hellerstein
 - Need to try different techniques based on the data
- Often the volume of data may be too much (e.g., internet routers)
 - Approximation techniques often used



Time Series Outliers

- Often data is in the form of a time series
- Rich literature on *forecasting* in time series data
- Can use the historical values / patterns in the data to flag outliers
 - Rolling standard deviation or MAD (median absolute variation)

Split-Apply-Combine

- The Split-Apply-Combine Strategy for Data Analysis, Wickam, 2011 ([here](#))
- Common data analysis pattern
 - Split: break a big problem into manageable pieces
 - Apply: operate on each piece independently
 - Combine: combine the pieces back together
- Pros
 - Code is compact
 - Easy to parallelize
- E.g.,
 - group-wise ranking
 - group vars (e.g., sums, means, counts)
 - create new models per group
- Supported by many languages
 - Pandas
 - SQL GROUP BY operator
 - Map-Reduce

```
In [94]: animals.groupby("kind").height.agg(  
.....:     min_height="min",  
.....:     max_height="max",  
.....: )  
.....:
```

```
Out[94]:
```

	min_height	max_height
kind		
cat	9.1	9.5
dog	6.0	34.0