

# UMD DATA605 - Big Data Systems

## Graph Data Management

### Neo4J

Dr. GP Saggese  
[gsaggese@umd.edu](mailto:gsaggese@umd.edu)

with thanks to  
Amol Deshpande, R. Licher (Technion),  
S. Nagarajan (U. Texas)

# Overview

- Motivation
- Graph data models
- Storing graph data
- Querying graph data
- Typical graph analysis tasks
- Executing graph analysis tasks

# Overview


- Motivation
- Graph data models
  - E.g., RDF, Property Graph, XML
- Storing graph data
  - E.g., Neo4j
- Querying graph data
  - E.g., Cypher, SPARQL, Gremlin
- Typical graph analysis tasks
  - E.g., PageRank, clustering
- Executing graph analysis tasks
  - E.g., Google Pregel, Apache Giraph, Spark GraphX

# Overview


- **Motivation**
- Graph data models
- Storing graph data
- Querying graph data
- Typical graph analysis tasks
- Executing graph analysis tasks

# Graphs: Background

- A *graph* (or *network*) captures a set of entities and interconnections between pairs of them
  - Entities / objects represented by *vertices* or *nodes*
  - Interconnections between pairs of vertices called *edges* (or *links*, *arcs*, *relationships*)
- Graph theory and algorithms widely studied in Computer Science
  - Not as much work on managing graph-structured data




*An undirected, unweighted graph*




*A directed, edge-weighted graph*

# Graph Data Structures: Motivation

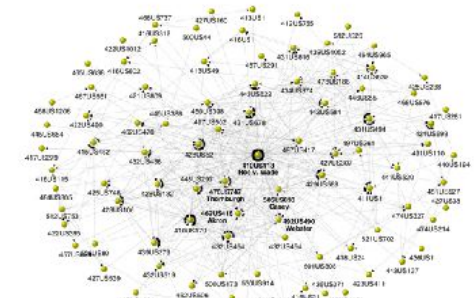
- Increasing interest in querying and reasoning about the *underlying graph structure* in a variety of disciplines




A protein-protein interaction network




Social networks



Supreme court citation network

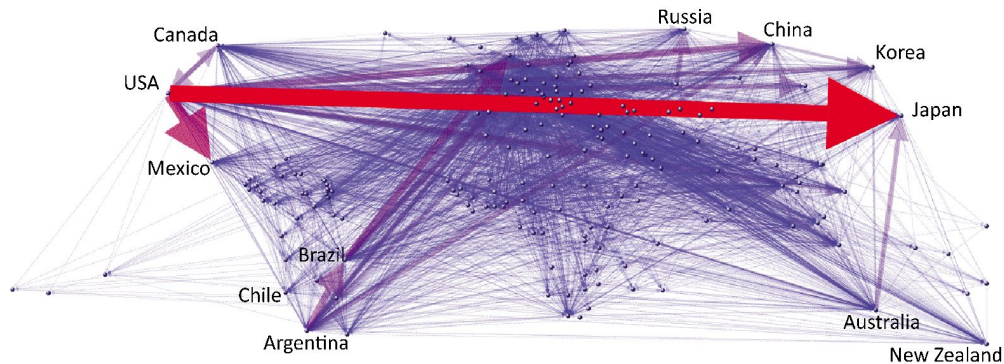


Financial transaction networks

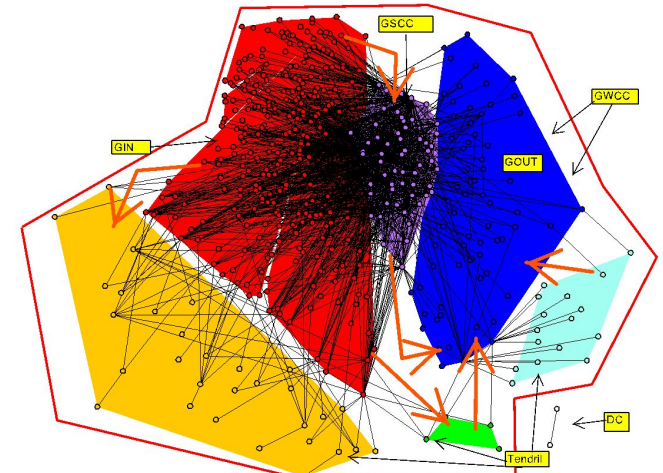


Stock Trading Networks

# Motivation



*Global virtual water trade network*



*Federal funds networks*

*Citation networks*

*Parcel shipment networks*

*Collaboration networks*

*Knowledge Graph*

*Telecommunications networks*

*World Wide Web*

*Disease transmission networks*

# Motivation

- **Graph data structures have not changed that much over time**
  - Same problems in representing the data in 1960s than today
- **What has changed in recent years**
  - Large data volumes and easier availability
  - Reasoning about the graph structure can provide useful and actionable insights
    - Lose too much information if graph structure ignored
  - Not easy to query using traditional tools (e.g., relational DBs)
    - Need specialized tools (e.g., Neo4j)
  - Hard to efficiently process graph-structured queries using existing tools
    - Dedicated solutions: Google Pregel / Apache Giraph, Spark GraphX
    - Problems getting worse with increasingly large graphs seen in practice



# Overview

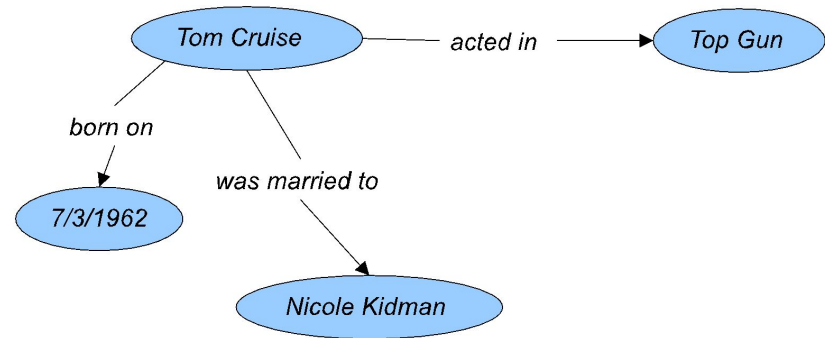
- Motivation
- **Graph data models**
- Storing graph data
- Querying graph data
- Typical graph analysis tasks
- Executing graph analysis tasks

# Knowledge Graphs

- **Representation of knowledge in the form of graphs**
  - Capture entities, relationships, and properties
  - Provide a structured view of real-world information
- Can be represented using RDF or Property Graph models
  - E.g., Google Knowledge Graph, DBpedia, Wikidata
- **Applications**
  - Enable machine understanding of complex domains
  - Support semantic search, recommendation, and analytics
  - Used in various industries for data integration, knowledge discovery, and AI applications
- **Ontologies**
  - Provide a formal representation of knowledge
  - Promote interoperability across knowledge bases

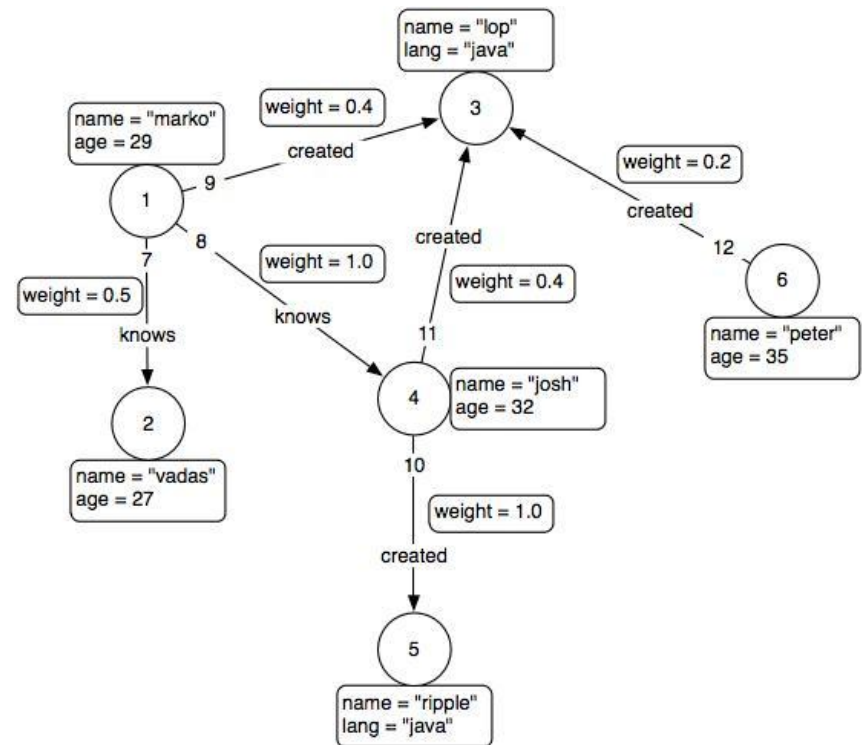

# Graph Data Models: RDF

- [Resource Description Framework](#)
- RDF uses triples subject-predicate-object
  - Each triple connects a “subject” and an “object” through a “predicate”
  - E.g., “TomCruise-acted-TopGun”
- **Used to represent *knowledge bases***
  - Typically queried through SPARQL
- **Pros**
  - Standardization
    - Standard W3C to model data
    - Subject and object can be URI (Uniform Resource Identifier) in semantic web
  - Interoperability
    - Can merge RDF data store
  - Extensibility
    - Can add new nodes and relationships
    - Support ontologies



# Graph Data Models: Property Graph

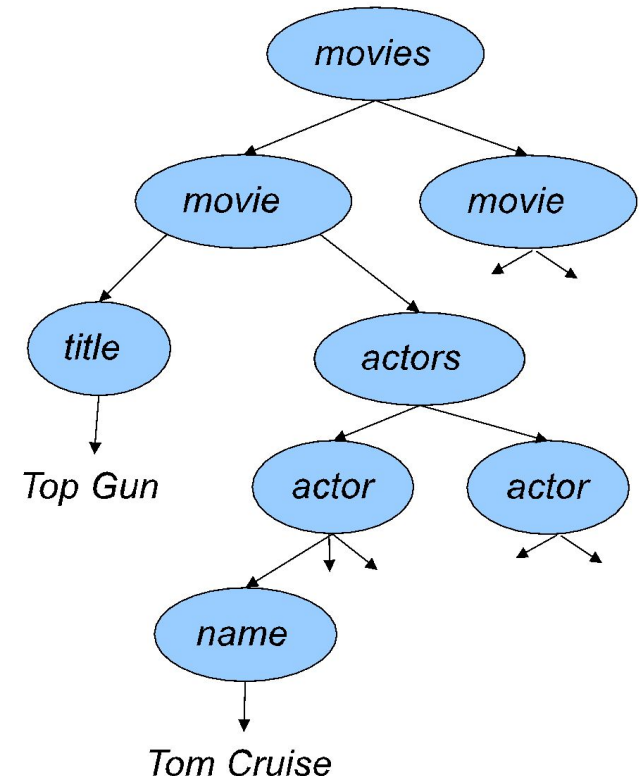
- A directed graph where each node and each edge may be associated with a set of *properties* (key-values)
- Query languages
  - Cypher (e.g., Neo4j)
  - Gremlin (e.g., Apache TinkerPop)
- Lack universal standard
- Similar expressive power to RDFs but less “schema” so more difficult to interoperate
- Used by many open-source graph data management tools



# Graph Data Models: XML

- Commonly used data model for representing data without rigid structure
- It is a directed labeled tree
- Popular data exchange format for non-tabular data

```
<movies>
  <movie>
    <title>Top Gun</title>
    <actors>
      <actor>
        <name>Tom Cruise</name>
        <born>ti/3/1962</born>
      </actor>
      <actor>
        ...
      </actor>
    </actors>
  </movie>
  ...
</movies>
```



# Overview

- Motivation
- Graph data models
- **Storing graph data**
- Querying graph data
- Typical graph analysis tasks
- Executing graph analysis tasks

# Storing Graph Data

1. **File systems**
  - + Very simple
  - No support for transactions, ACID
  - Minimal functionality (e.g., must build the analysis/querying on top)
2. **Relational database**
  - + Mature technology
  - + All the good stuff (SQL, transactions, ACID, toolchains)
  - Minimal functionality
3. **NoSQL key-value stores**
  - + Can handle very large datasets efficiently in a distributed fashion
  - Minimal functionality
4. **Graph database**
  - + Efficiently support for queries / tasks (e.g., graph traversals)
  - Not as mature as RDBMs
  - Often no declarative language (similar to SQL)
    - You need to write programs

# Graph Databases

- Many specialized [graph database systems](#) in recent years
  - E.g., Neo4j, Titan, OrientDB, AllegroGraph
- A few key distinctions from relational databases
  - Built to manage and query graph-structured data
  - Store the graph structure explicitly using data structures with pointers
    - Avoid the need for joins, making graph traversals easier
    - More natural to write *queries* and *graph algorithms* (reachability or shortest paths)
  - Support graph query languages like SPARQL, Cypher, Gremlin
  - Fairly rudimentary declarative interfaces
  - Most applications need to be written using programmatic interfaces
  - Expose a programmatic API to write arbitrary graph algorithms



# Overview

- Motivation
- Graph data models
- Storing graph data
- **Querying graph data**
- Typical graph analysis tasks
- Executing graph analysis tasks

# Query Languages for Graph Databases

- [Cypher](#)

- Designed for Property Graphs
  - Data = vertices and edges annotated with key-value properties
- Declarative
- Subgraph pattern matching
- Can't easily handle queries like reachability
- Native query language for Neo4j

- [Gremlin](#)

- Works with both RDF and Property Graphs
- Imperative
- Allow to describe graph traversal

- [SPARQL](#)

- Similar to Cypher
- Query language for RDF data
- Standardized by W3C

```
MATCH (nicole:Actor
      {name: 'Nicole Kidman'})-[:ACTED_IN]->(movie:Movie)
WHERE movie.year < 2007
RETURN movie
```


```
// calculate basic collaborative filtering for vertex 1
m = [:]
g.v(1).out('likes').in('likes').out('likes').groupCount(m)
m.sort{-it.value}

// calculate the primary eigenvector (eigenvector centrality) of a graph
m = [:]; c = 0;
g.V.as('x').out.groupCount(m).loop('x'){c++ < 1000}
m.sort{-it.value}
```

```
PREFIX foaf: <http://xmlns.com/foaf/0.1/>
SELECT ?name
       ?email
WHERE
{
  ?person a          foaf:Person .
  ?person foaf:name   ?name .
  ?person foaf:mbox    ?email .
}
```

# Neo4j

- Graph DB storing data as Property Graph
  - Nodes, edges hold data as key-value pairs (like non-relational DBs)
- Focus is
  - On relationships between values
  - Instead of commonalities among sets of values (e.g., tables of rows for RDBMs and collection of documents)
- Two querying languages
  - Cypher, Gremlin
- GUI or REST API
- Full ACID-compliant transactions (atomicity, consistency, isolation, durability)
- High-availability clustering
- Incremental backups
- Can run in small application or run on large clusters of servers



# Graph DB: Example

## • Specs

- Create a wine suggestion engine
- Wines are categorized by different
  - Varieties (e.g., Chardonnay, Pinot Noir)
  - Regions (e.g., Bordeaux, Napa, Tuscany)
  - Vintage (year in which the grapes were harvested)
- Keep track of articles describing wines by various authors
- Users can track their favorite wines

## • Relational model


- The important relationships are `produced`, `reported\_on`, `grape\_type`
- Create various tables
  - `wines`: (id, name, year)
  - `wines\_categories` (wine\_id, category\_id)
  - `category` table (id, name)
  - `wines\_articles` (wine\_id, article\_id)
  - `articles` (id, publish\_date, title, content)

## • Problem with relational approach

- There isn't much of a schema
- Lots of incomplete data
- An old saying in relational DB world: *"On a long enough timeline all fields become optional"*

## • Graph DB approach

- Provide values and structure only where necessary



# Labeled Property Graphs in Neo4j

- **Nodes**

- Main data elements
- Connected to other nodes via *relationships*
- Can have one or more *properties* (stored as key/value pairs)

- **Relationships**

- Connect two *nodes*
- Are directional
- *Nodes* can have multiple relationships
- Can have one or more *properties* (stored as key/value pairs)

- **Properties**

- Named values where the name (or key) is a string
- Can be indexed and constrained
- Composite indexes can be created from multiple properties

- **Labels**

- Used to group nodes into sets
- A node may have multiple labels
- Labels indexed to accelerate finding nodes in the graph
- Native label indexes optimized for performance

# Cypher Example

# Create a wine node with attributes.

```
$ CREATE (w:Wine
  {name:"Prancing Wolf",
   style: "ice wine",
   vintage: 2015})
```

# Return the entire graph.

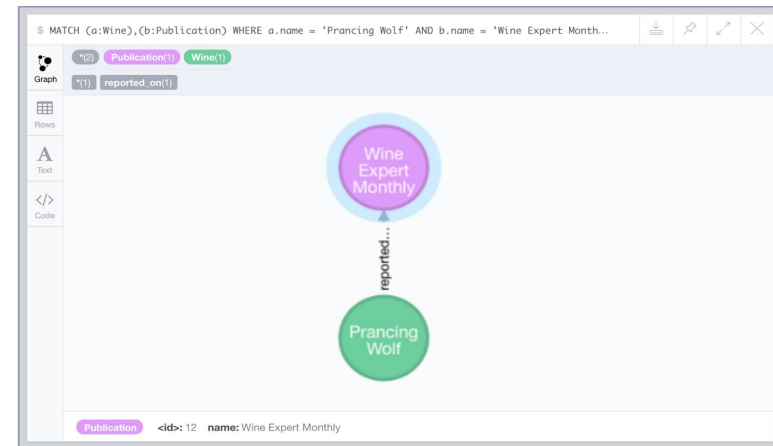
```
$ MATCH (n)
RETURN n;
```

# Create a publication node.

```
$ CREATE (p:Publication
  {name: "Wine Expert Monthly"})
```

# Create a relation “reported\_on”.

```
$ MATCH (p:Publication
  {name: "Wine Expert Monthly"}),
(w:Wine {name: "Prancing Wolf",
  vintage: 2015})
CREATE (p)-[r:reported_on]->(w)
```



# Cypher Example

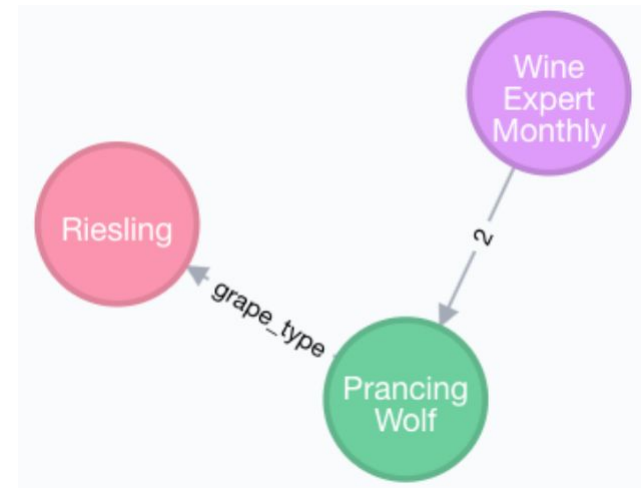
# Attach a rating.

```
$ MATCH (p:Publication {name: "Wine Expert Monthly"}),  
      (w:Wine {name: "Prancing Wolf"})  
  CREATE (p)-[r:reported_on {rating: 2}]->(w)
```

# Add a “grape\_type” relationship.

```
$ CREATE (g:GrapeType {name: "Riesling"})
```

```
$ MATCH (w:Wine {name: "Prancing Wolf"}),  
      (g:GrapeType {name: "Riesling"})  
  CREATE (w)-[r:grape_type]->(g)
```



# Cypher Example

```
# Add winery.
$ CREATE (wr:Winery {name: "Prancing Wolf Winery"})

# Add "produced" relationship.
$ MATCH (w:Wine {name: "Prancing Wolf"}),
      (wr:Winery {name: "Prancing Wolf Winery"})
  CREATE (wr)-[r:produced]->(w)

$ CREATE (w:Wine
  {name:"Prancing Wolf", style: "Kabinett", vintage: 2002})
$ CREATE (w:Wine
  {name: "Prancing Wolf", style: "Spätlese", vintage: 2010})
$ MATCH (wr:Winery
  {name: "Prancing Wolf"}),(w:Wine {name: "Prancing Wolf"})
  CREATE (wr)-[r:produced]->(w)

# Add "grape type" relationship.
$ MATCH (w:Wine), (g:GrapeType {name: "Riesling"})
  CREATE (w)-[r:grape_type]->(g)
```





# Cypher Example

- Add a social component to the wine graph
  - People preference for wine
  - Relationships with one another

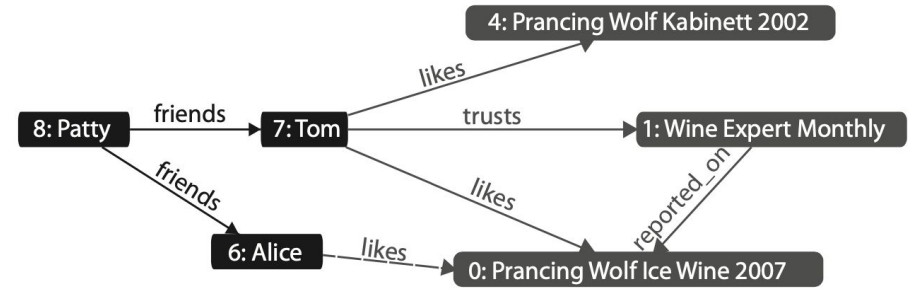
# Alice likes a certain wine.

```
$ CREATE (p:Person {name: "Alice"})
$ MATCH (p:Person {name: "Alice"}),
      (w:Wine {name: "Prancing Wolf",
               style: "ice wine"})
  CREATE (p)-[r:likes]->(w)
```

# Patty and Tom are friends.

```
$ CREATE (p:Person {name: "Patty"})
$ MATCH (p1:Person {name: "Patty"}),
      (p2:Person {name: "Tom"})
  CREATE (p1)-[r:friends]->(p2)
```

- The changes were made  
“superimposing” new relationships  
without changing the previous data



# Cypher Example

# See all nodes associated with Alice.


```
$ MATCH (p:Person
{name: "Alice"})-->(n)
RETURN n;
```

# Find all of the people that Alice is friends with, returning only the name property of those nodes

```
$ MATCH (p:Person
{name: "Alice"})-->(other: Person)
RETURN other.name;
```

# Find friends of friends of Alice.

```
$ MATCH
(fof:Person)-[:friends]-(f:Person)-[:f
riends]-(p:Person {name: "Patty"})
RETURN fof.name;
```



# A general query structure

```
MATCH [Nodes and relationships]
WHERE [Boolean filter statement]
RETURN [DISTINCT] [statements [AS alias]]
ORDER BY [Properties] [ASC\DESC]
SKIP [Number] LIMIT [Number]
```

# Simple query

Get all nodes of type *Program* that have the name *Hello World!*

```
MATCH (a : Program)
WHERE a.name = 'Hello World!'
RETURN a
```



# Query relationships

Get all relationships of type *Author* connecting *Programmers* and *Programs*:



```
MATCH (a:Programmer)-[r:Author]->(b:Program)
RETURN r
```

# Matching nodes and relationships

- Nodes

`(a), (), (:Ntype), (a:Ntype),  
(a { prop:'value' } ) ,  
(a:Ntype { prop:'value' } )`

- Relationships

`(a)--(b)  
(a)-->(b), (a)<--(b),  
(a)-->(), (a)-[r]->(b),  
(a)-[:Rtype]->(b), (a)-[:R1|:R2]->(b),  
(a)-[r:Rtype]->(b)`

- May have more than 2 nodes

`(a)-->(b)<--(c), (a)-->(b)-->(c)`

- Path

`p = (a)-->(b)`

# More options

- Relationship distance:
  - $(a) - [:Rtype^*2] \rightarrow (b)$  – 2 hops of type Rtype.
  - $(a) - [:Rtype^*] \rightarrow (b)$  – any number of hops of type Rtype.
  - $(a) - [:Rtype^*2..10] \rightarrow (b)$  – 2-10 hops of Rtype.
  - $(a) - [:Rtype^* ..10] \rightarrow (b)$  – 1-10 hops of Rtype.
  - $(a) - [:Rtype^*2.. ] \rightarrow (b)$  – at least 2 hops of Rtype.

Could be used also as:

$(a) - [r^*2] \rightarrow (b)$  – r gets a sequence of relationships  
 $(a) - [* \{prop:val\}] \rightarrow (b)$

# Operators

- Mathematical  
+, -, \*, /, %, ^ (power, not XOR)
- Comparison  
=, <>, <, >, >=, <=, =~ (Regex), IS NULL ,  
IS NOT NULL
- Boolean  
AND, OR, XOR, NOT
- String  
Concatenation through +
- Collection  
Concatenation through +  
IN to check if an element exists in a collection



# More WHERE options

- WHERE others.name IN ['Andres', 'Peter']
- WHERE user.age IN range (18,30)
- WHERE n.name =~ 'Tob.\*'
- WHERE n.name =~ '(?i)ANDR.\*' - (case insensitive)
- WHERE (tobias)-->()
- WHERE NOT (tobias)-->()
- WHERE has(b.name)
- WHERE b.name? = 'Bob'

(Returns all nodes where name = 'Bob' plus all nodes without a name property)

# Functions

- On paths:
  - MATCH shortestPath( (a)-[\*]-(b) )
  - MATCH allShorestPath( (a)-[\*]-(b) )
  - Length(path) – The path length or 0 if not exists.
  - RETURN relationships(p) - Returns all relationships in a path.
- On collections:
  - RETURN a.array, filter(x IN a.array WHERE length(x)= 3)  
    **FILTER** - returns the elements in a collection that comply to a predicate.
  - WHERE ANY (x IN a.array WHERE x = "one" ) – at least one
  - WHERE ALL (x IN nodes(p) WHERE x.age > 30) – all elements
  - WHERE SINGLE (x IN nodes(p) WHERE var.eyes = "blue") – Only one
  - \* nodes(p) – nodes of the path p

# With

- Manipulate the result sequence before it is passed on to the following query parts.
- Usage of WITH:
  - Limit the number of entries that are then passed on to other MATCH clauses.
  - Introduce aggregates which can then be used in predicates in WHERE.
  - Separate reading from updating of the graph. Every part of a query must be either read-only or write-only.

# Data access is *programmatic*

- REST API
- Through the Java APIs
  - JVM languages have bindings to the same APIs
    - JRuby, Jython, Clojure, Scala...
- Managing nodes and relationships
- Indexing
- Traversing
- Path finding
- Pattern matching

# Overview

- Motivation
- Graph data models
- Storing graph data
- Querying graph data
- **Typical graph analysis tasks**
- Executing graph analysis tasks

# Queries vs Analysis Tasks

## . Queries

- Focused exploration of the data
- Result is typically a small portion of the graph (often just a node)
- Challenges
  - Minimize the portion of the graph that is explored
  - Use of indexes (auxiliary data structures)

## . Analysis tasks


- Typically require processing the entire graph
- Challenges
  - How to handle the large volume of data efficiently
  - How to parallelize if data doesn't fit in memory / disk

# Examples of Graph Queries / Tasks

- Subgraph pattern matching
  - Find matching instances of a given small graph in a large graph
  - Although technically NP-hard, usually the patterns are small
- Shortest path queries
  - Find the shortest path between two given nodes
  - E.g., in road networks
- Reachability
  - Given two nodes, is there an undirected or directed path between them?
  - Sometimes with constraints on the types of edges that can be used
- Keyword search
  - Find the smallest subgraph that contains all the specified keywords
- Historical queries
  - Given a node, find other nodes that evolved most similarly in the past
- Graph algorithms
  - Network flows
  - Spanning trees

# Queries: Subgraph Matching


- Given a “query” graph, find where it occurs in a given “data” graph
  - Query graph can specify restrictions on the graph structure, on values of node attributes, and so on
  - An important variation: *approximate* matching





# Queries: Connection Subgraphs


- Given a data graph and two (or more) nodes in it, find a small subgraph that best captures the relationship between the nodes
- How to define “best captures”?
  - E.g., “shortest path”: but that may not be most informative



*The “red” path between D and J  
maybe more informative than the  
“green” path*


# Graph Analysis: Centrality Measures

- Centrality measure: a measure of the relative importance of a vertex within a graph
- Many different definition of centrality measures
  - Can give fairly different results
- **Degree centrality of a node  $u$** 
  - Number of edges incident on  $u$
- **Betweenness centrality of a node  $u$** 
  - Number of shortest paths between pairs of vertices that go through  $u$
- **Pagerank of a node  $u$ :**
  - Probability that a random surfer (who is following links randomly) ends up at node  $u$



# Graph Analysis: Community Detection

- Goal: partitioning the vertices into (potentially overlapping) groups based on the interconnections between them
  - Basic intuition: More connections within a community than across communities
  - Provide insights into how networks function; identify functional modules; improve performance of Web services; etc.
- Numerous techniques proposed for community detection over the years
  - Graph partitioning-based methods
  - Maximizing some “goodness” function
  - Recursively removing high centrality edges
  - And so on ...




# Overview

- Motivation
- Graph data models
- Storing graph data
- Querying graph data
- Typical graph analysis tasks
- **Executing graph analysis tasks**

# Bulk Synchronous Parallel (BSP)

- BSP model is a computational model used to design parallel algorithms for distributed systems
- Computation is divided into a series of *supersteps*, each consisting of three main phases
  - **Local computation phase**
    - Each processing unit performs calculations independently and concurrently, without any interaction
  - **Communication phase**
    - The processing units exchange information with each other by sending and receiving messages
    - These messages can be exchanged asynchronously without waiting for a response
  - **Synchronization phase**
    - Aka barrier
    - Ensures that all processing units have completed their local computations and communication before proceeding to the next superstep
    - This guarantees that all messages from the previous superstep have been received and processed
- Suitable for iterative graph algorithms
  - E.g., PageRank and Shortest Path

# Bulk Synchronous Parallel (BSP)



# Pregel System

- Large-scale graph processing system developed by Google
  - [Pregel paper](#), 2010
- Inspired by the Bulk Synchronous Parallel (BSP) model
  - Vertex-centric programming model
  - Asynchronous message passing between vertices
- Fault-tolerant using checkpointing mechanism
- Scalable and distributed architecture
- Designed for processing large graphs with billions of vertices and edges
- Handles graph mutations and updates during computation
- Not open-source, used internally at Google

# Apache Giraph



- [Apache Giraph](#)
- Open-source graph processing framework, inspired by Google's Pregel
- Implemented by Facebook and then open-sourced
- Built on top of Apache Hadoop
- Fault-tolerant using Hadoop's checkpointing mechanism
- Scalable and distributed architecture
- Suitable for large-scale graph analytics and machine learning algorithms
- Actively maintained and widely adopted in the open-source community



# Apache Spark GraphX

- [Apache Spark GraphX](#)
- Graph processing library for Apache Spark
- Built on top of Spark's RDD (Resilient Distributed Dataset) model
- Supports both directed and undirected graphs
- Provides a flexible graph computation API
- Optimized for iterative graph computations
- Scalable and fault-tolerant architecture
- Supports in-memory graph processing for improved performance
- Suitable for large-scale graph analytics and machine learning tasks
- Implements various graph algorithms
  - E.g., PageRank, Connected Components, and Shortest Path

